


Introduction to Amazon EKS

Outline

- **Kubernetes Basics**
- **EKS Basics**
- **Logging And Monitoring**
- **EKS Advanced Concepts**
- **Securing EKS**
- **Fargate**
- **Deploying EKS with DevOps**
- **Real World EKS Projects**

What is EKS?

Kubernetes Architecture



Kubernetes Control Plane - Self Managed



Amazon EC2

- Need to make Control Plane Highly Available
 - Maintain multiple EC2 in multiple AZ
- Scale Control Plane if needed
- Keep etcd up and running
- Overhead of managing EC2s
 - AMI Rehydration
 - Security Patching
 - Replace failed EC2s
 - Orchestration for Kubernetes Version Upgrade

Kubernetes Control Plane - AWS Managed



AWS Manages Kubernetes Control Plane

- AWS maintains High Availability - Multiple EC2s in Multiple AZs
- AWS Detects and Replaces Unhealthy Control Plane Instances
- AWS Scales Control Plane
- AWS Maintain etcd
- Provides Automated Version Upgrade and Patching
- Supports Native and Upstream Kubernetes
- Integrated with AWS Ecosystem

EKS Data Plane



Amazon EC2

Self Managed Node Groups

- You maintain worker EC2s
- You orchestrate version upgrade, security patching, AMI Rehydration, keeping pods up during upgrade
- Can use custom AMI



Amazon EC2

Managed Node Groups

- AWS manages worker EC2s
- AWS provides AMI with security patches, version upgrade
- AWS manages pod disruption during upgrade
- Doesn't work with custom AMI



AWS Fargate

- No worker EC2 whatsoever!
- You define and deploy pods
- Container + Serverless!

EKS in AWS Ecosystem



AWS Secrets Manager



Amazon API Gateway



Elastic Load Balancing



Amazon EKS



AWS Identity and Access Management (IAM)



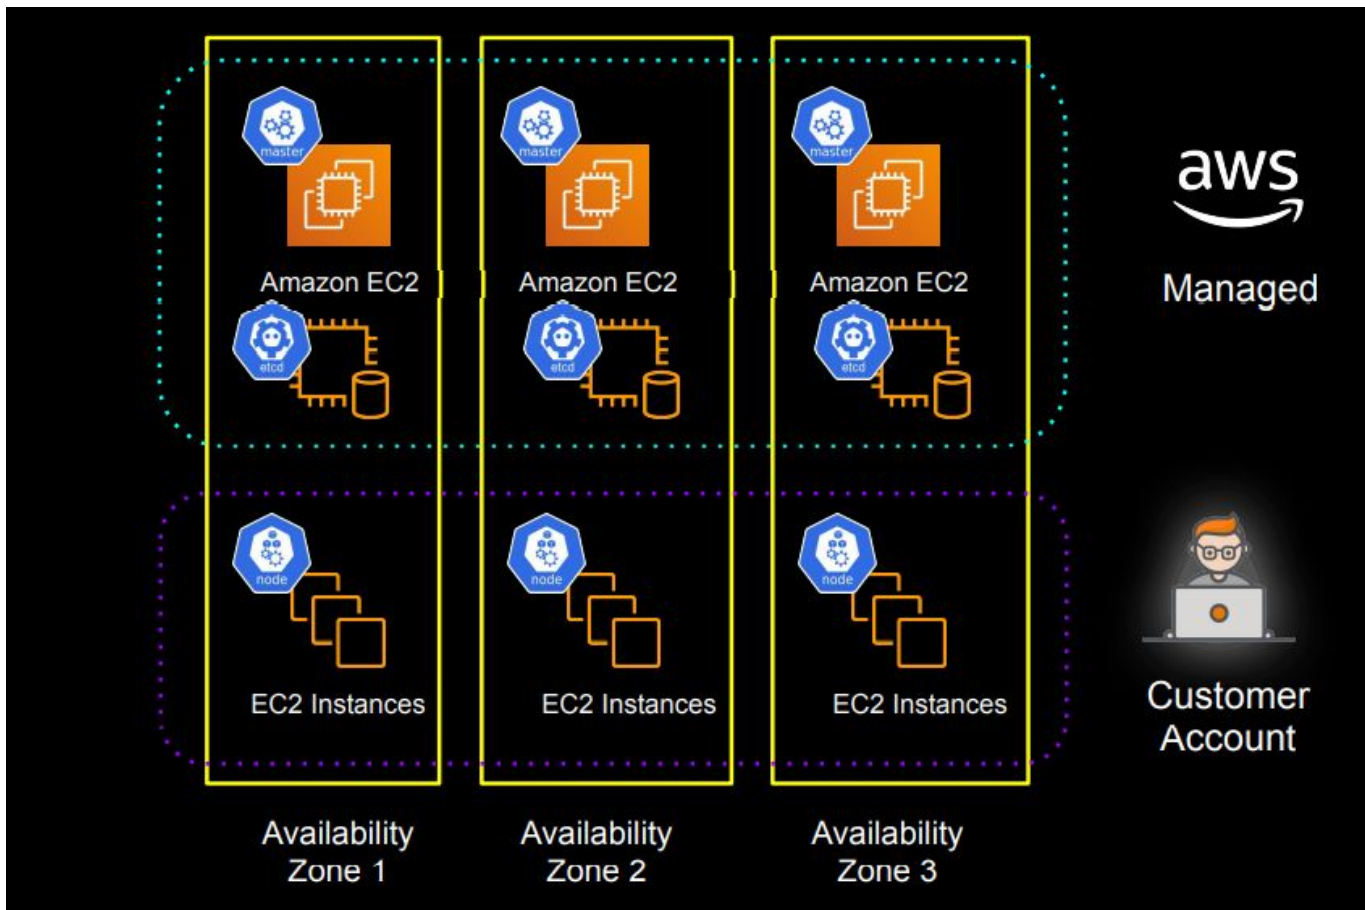
AWS CodePipeline



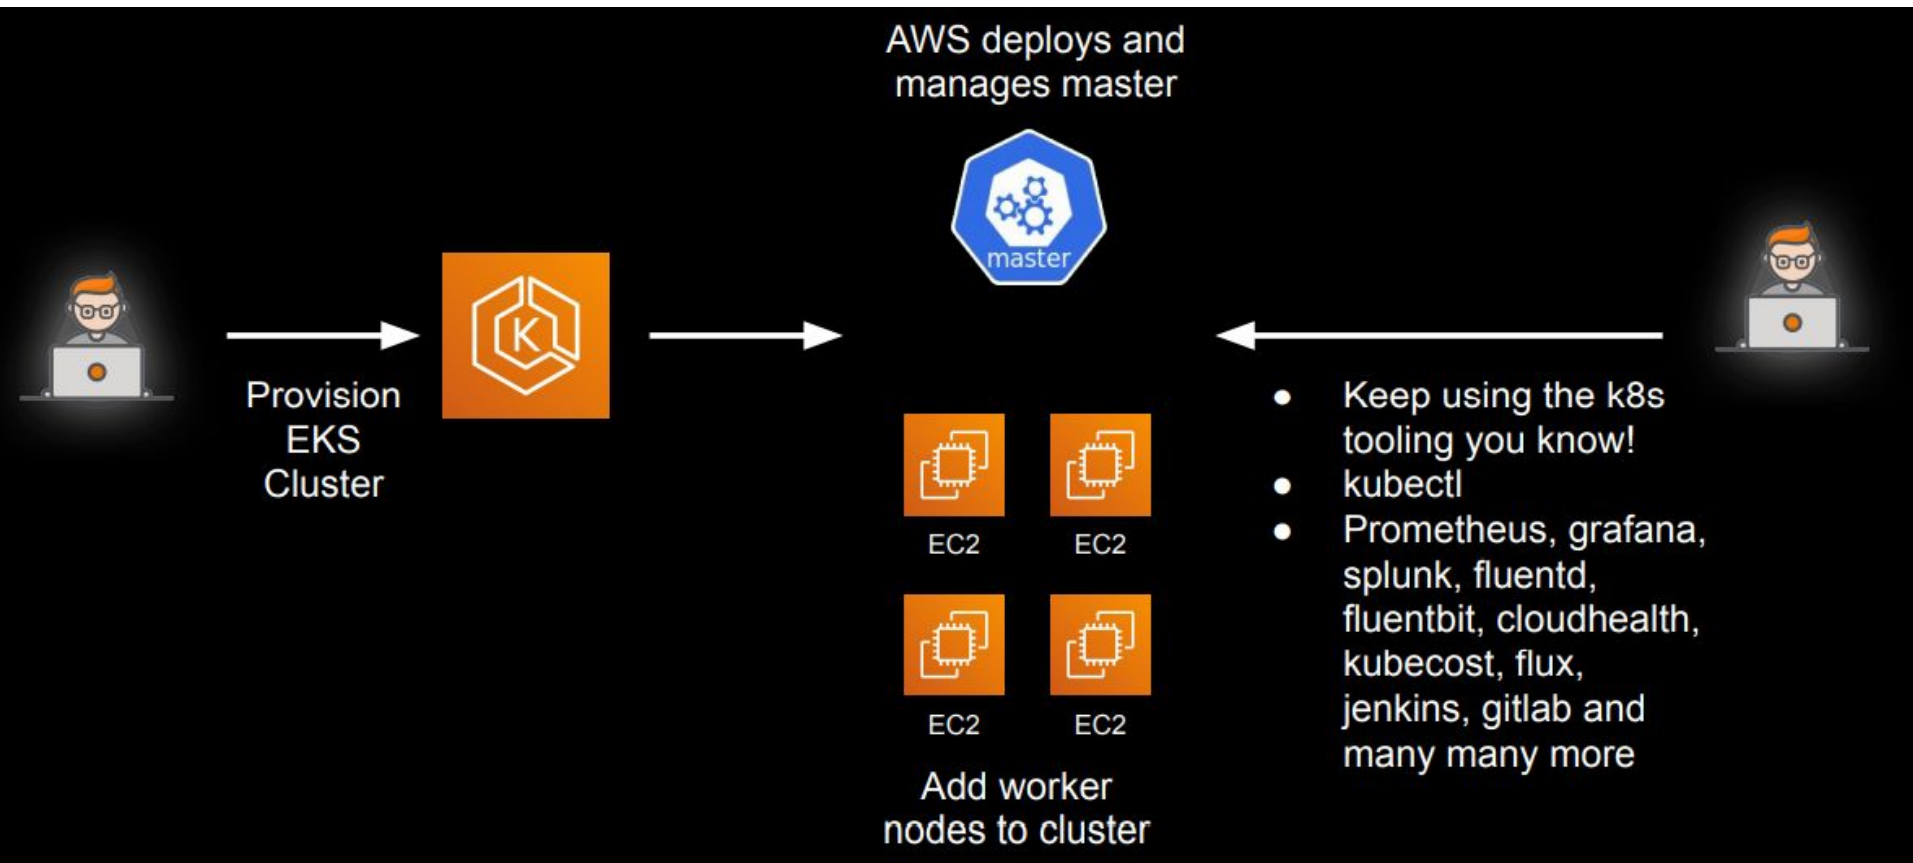
Amazon CloudWatch

And More....

EKS High Level Architecture



High Level Flow



Money Matters

Amazon Elastic Container Service for Kubernetes



Flat Charge 10 cents/Hour = \$72/Month



EC2



EC2



EC2

Price based on EC2

$$\begin{aligned}\text{EKS with 3 m5.large in us-east-1} &= \text{EKS Control Plane} + \text{Worker Nodes (m5.large X 3)} \\ &= \$72/\text{mo} \quad + \quad \$219.24/\text{mo} \\ &= \mathbf{\$291.24/\text{mo}}\end{aligned}$$

eksctl & kubectl

Ways To Spin Up EKS Cluster



AWS Console

CloudFormation

AWS CLI

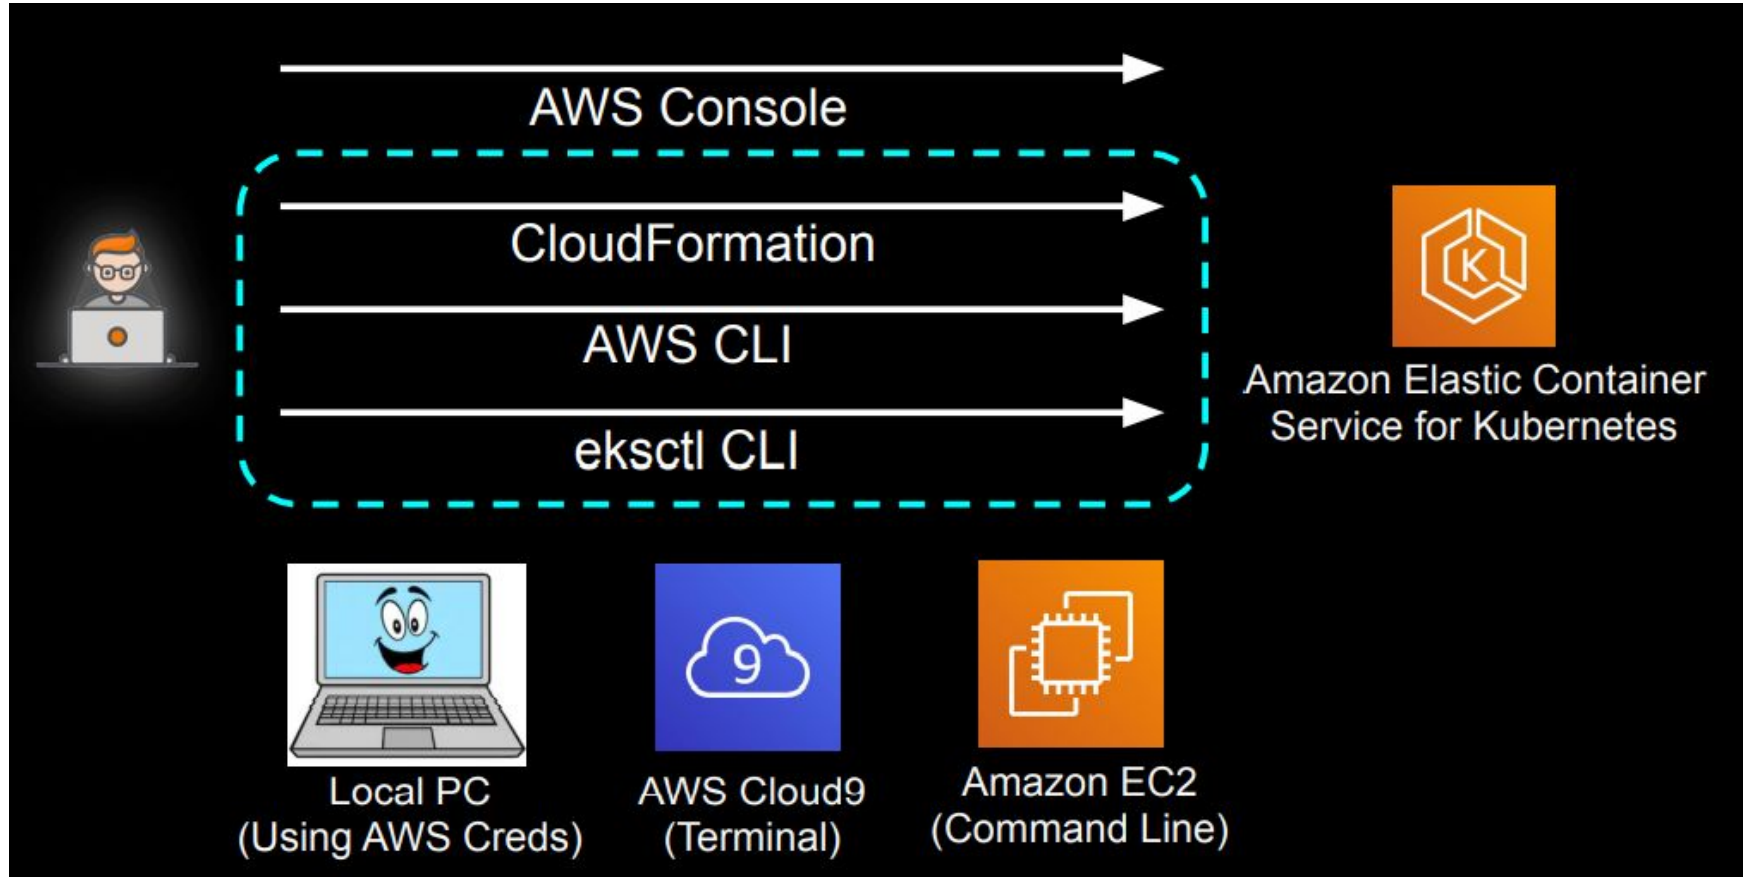


eksctl CLI

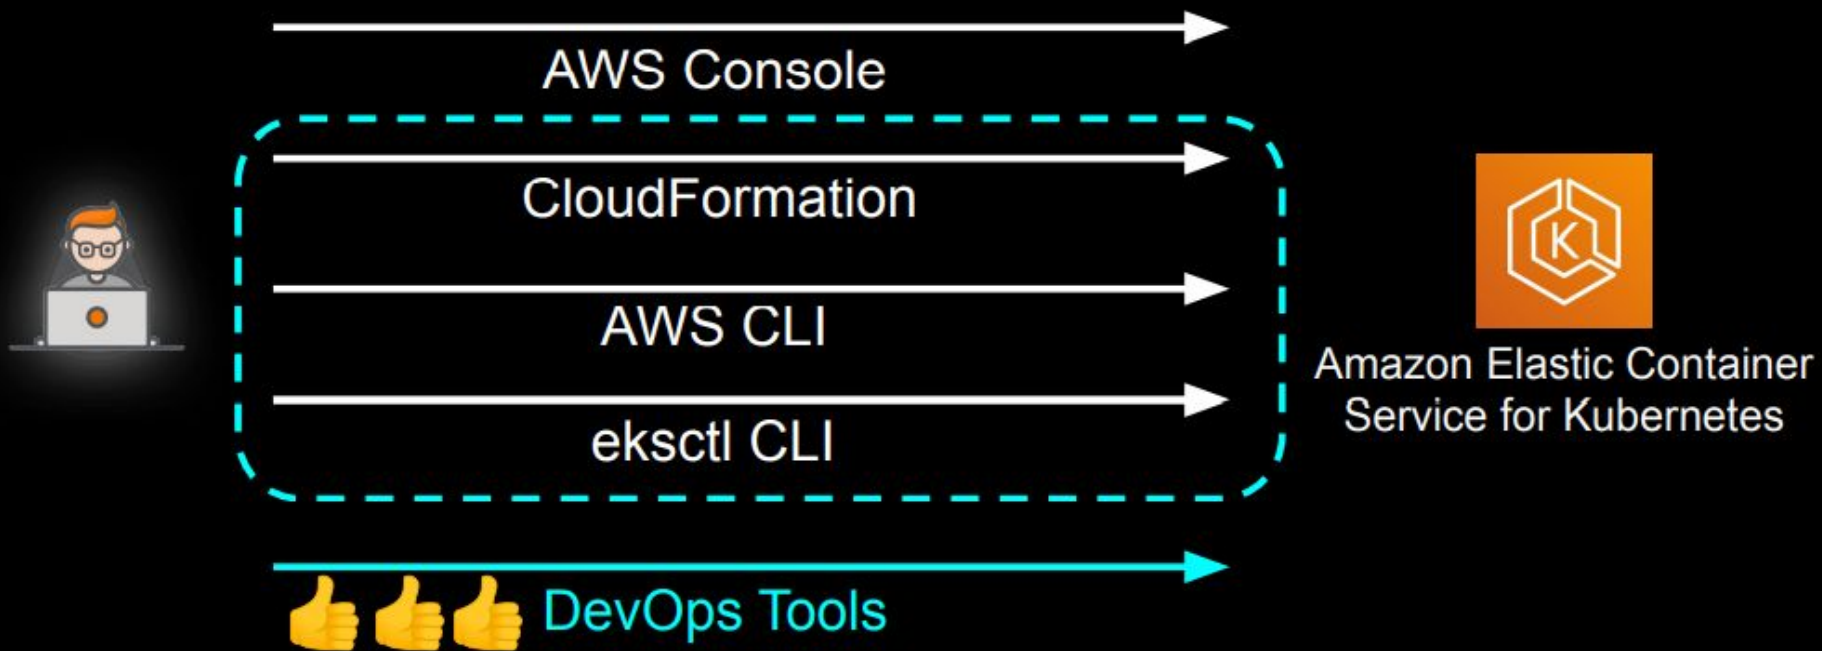


Amazon Elastic Container
Service for Kubernetes

Learning Medium



Automation



What is eksctl?

- CLI tool for creating clusters on EKS
- Easier than console, for real!
- Abstracts lots of stuff - VPC, Subnet, Sec. Group etc. using CloudFormation



eksctl create cluster



Amazon Elastic Container
Service for Kubernetes



AWS Fargate (On EKS)

Available eksctl features (Only on EKS)


- Create, drain and delete nodegroups
- Scale a nodegroup
- Update a cluster
- Use custom AMIs
- Configure VPC Networking
- Configure access to API endpoints
- Support for GPU nodegroups
- IAM Management and Add-on Policies
- Spot instances and mixed instances
- List cluster Cloudformation stacks
- Install coredns
- Write kubeconfig file for a cluster
- Spot instances and mixed instances
- Create, get, list and delete clusters

eksctl Commands

Command	Brief Description
<code>eksctl create cluster</code>	Create EKS Cluster with one nodegroup containing 2 m5.large
<code>eksctl create cluster --name <name> --version 1.15 --node-type t3.micro --nodes 2</code>	Create EKS Cluster with K8 version 1.15 with 2 t3.micro nodes
<code>eksctl create cluster --name <name> --version 1.15 --nodegroup-name <nodegrpname> --node-type t3.micro --nodes 2 --managed</code>	Create EKS cluster with managed node group
<code>eksctl create cluster --name <name> --fargate</code>	EKS Cluster with Fargate Profile

What is kubectl?

- CLI for running commands against a cluster on K8s resources
- Communicate via cluster API Server
- Works for any K8s cluster - EKS, K8s on EC2, GKE etc.



kubectl Command Syntax

kubectl

[command]

create
get
describe
delete
apply
attach
autoscale

And Many
More...

[type]

Any K8 Resource Type
pods OR pod OR po
namespaces OR ns
deployments OR deploy
replicasets OR rs

And Many More....

[name]

Optional
Name of the resource.
If omitted, all resource
details displayed

[flags]

Optional
--filename OR -f
--output OR -o

And Many More...

All available command and resource type : <https://kubernetes.io/docs/reference/kubectl/overview>

kubectl Command Syntax

kubectl	[command]	[TYPE]	[NAME]	[flags]
kubectl	get	pod	pod1	
kubectl	get	pod		
kubectl get po				
kubectl	get	pod	pod2	-o yaml




Most Used kubectl Commands

Command	Brief Description
<code>kubectl apply -f ./manifest-file.yaml</code>	Create resources based on manifest. Declarative Way! Best Way!
<code>kubectl get nodes</code>	List all node info
<code>kubectl get services</code>	List all services
<code>kubectl get pods -o wide</code>	List pods with more details
<code>kubectl get pod my-pod -o yaml</code>	Get a pod's YAML
<code>kubectl get deployment my-dep</code>	List a particular deployment
<code>kubectl exec -it podname -- /bin/bash</code>	Get a shell to the running Container

DEMO TIME!


- **Spin up EKS Cluster using eksctl**
- **Use kubectl**
 - **Deploy nginx using manifest file**
 - **Get resources info**

EKS Managed Nodegroups



EKS Regular Nodes and Pods

- K8 Version needs to be updated
- Deploy Security Patches
- EC2 AMI Need to Be Updated



EKS Nodes Regular Update

- K8 Version needs to be updated
- Deploy Security Patches
- EC2 AMI Need to Be Updated
 - You need to create AMIs



EC2 AMI




- Application will be DOWN!!
- Or You have to orchestrate HA

EKS Managed Nodegroups

- Create and Manage EC2 Workers for you
- Amazon releases AMIs with bug fix, security patches for EKS Worker Nodes
 - Custom AMIs not supported (As of April, 2020)
- Automated deployment of updated AMIs with security patches, CVEs
 - No app downtime
 - No overhead of user managed orchestration
 - Auto scaling group is used behind the scenes

EKS Nodes: Managed Nodegroup Update

- Worker Node K8s Version updated with one click/API call
- AWS Provides AMI with security updates
 - No patching/Rehydrate effort for you



- Application will be up and running
- AWS orchestrate HA through Auto Scaling Group
- Respects Pod Disruption Budget

EKS Logging & Monitoring

EKS Logging

**EKS Control Plane
Logging**

**EKS Worker Nodes
Logging**

EKS Logging

- **EKS Control Plane Logging**
 - **K8 api**
 - **audit**
 - **authenticator**
 - **controllerManager**
 - **scheduler**
- **EKS Worker Nodes Logging**

EKS Logging


- **EKS Control Plane Logging**

- **K8 api**
- **audit**
- **authenticator**
- **controllerManager**
- **scheduler**


- **EKS Worker Nodes Logging**

- **System logs from kubelet, kube-proxy, or dockerd**
- **Application logs from application containers**

Logging Caveat




- If instance terminated, logs are gone




- Logs need to be aggregated in a meaningful way

Logging Caveat

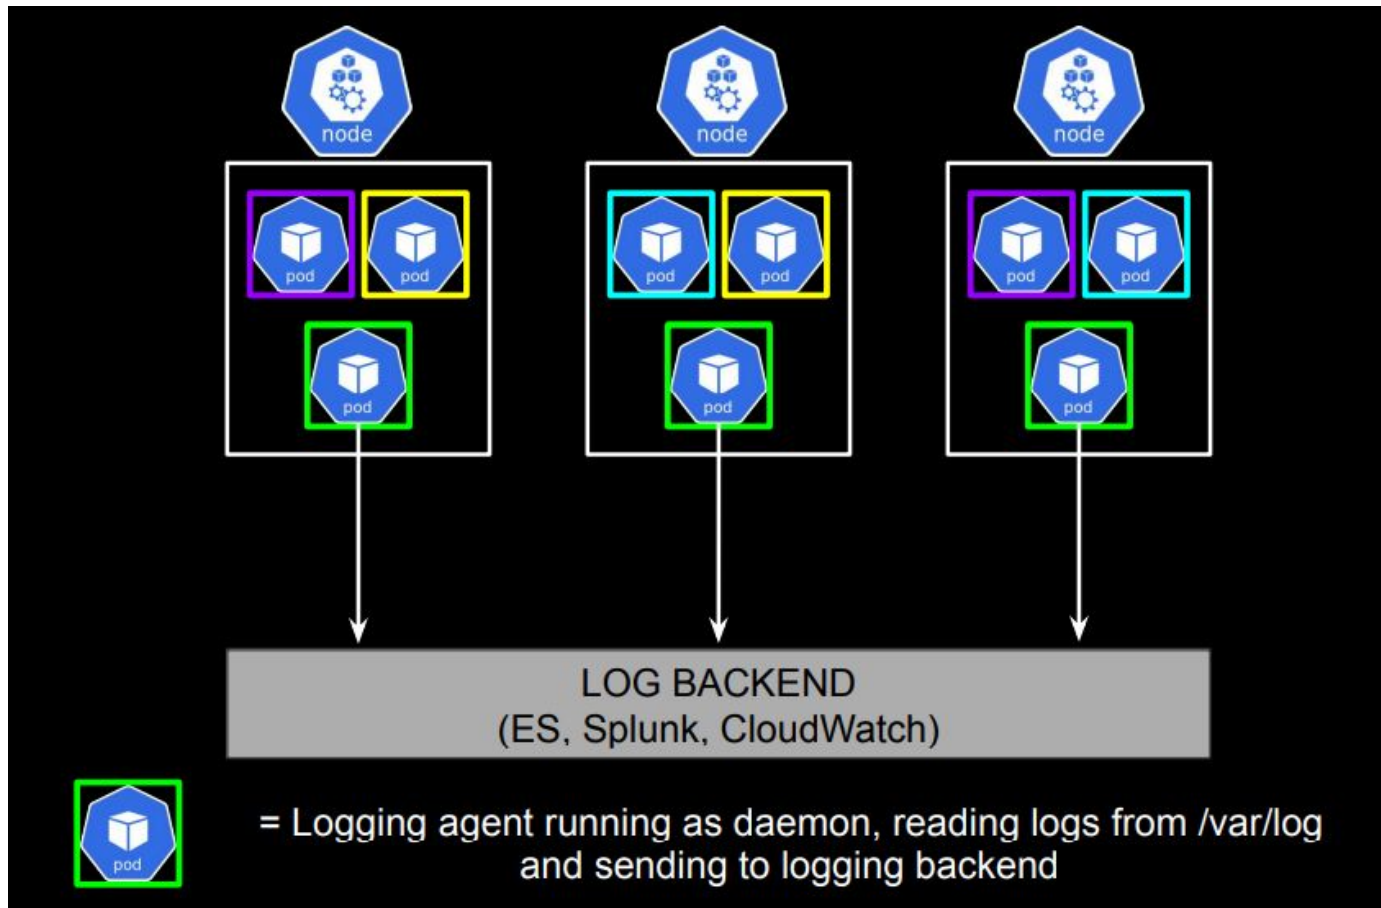


Kubernetes Worker Nodes

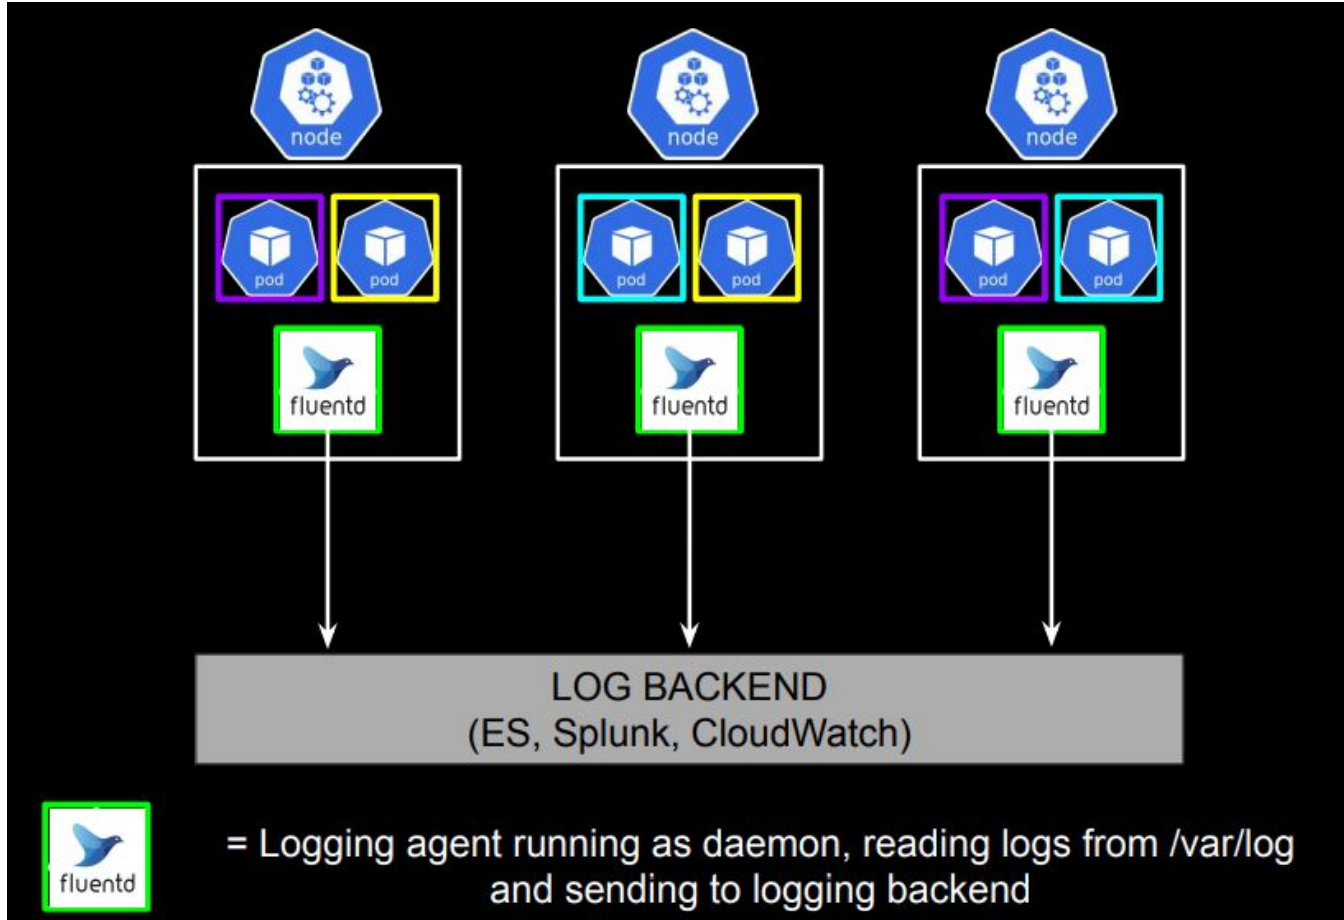


- Containerized application writes to
 - stdout and stderr
- System logs go to
 - systemd
- Container redirect logs to `/var/log/containers/*.log` files
- Now we know where to extract logs from!

Implementation



Implementation



Different Options

Agent



Logging Backend

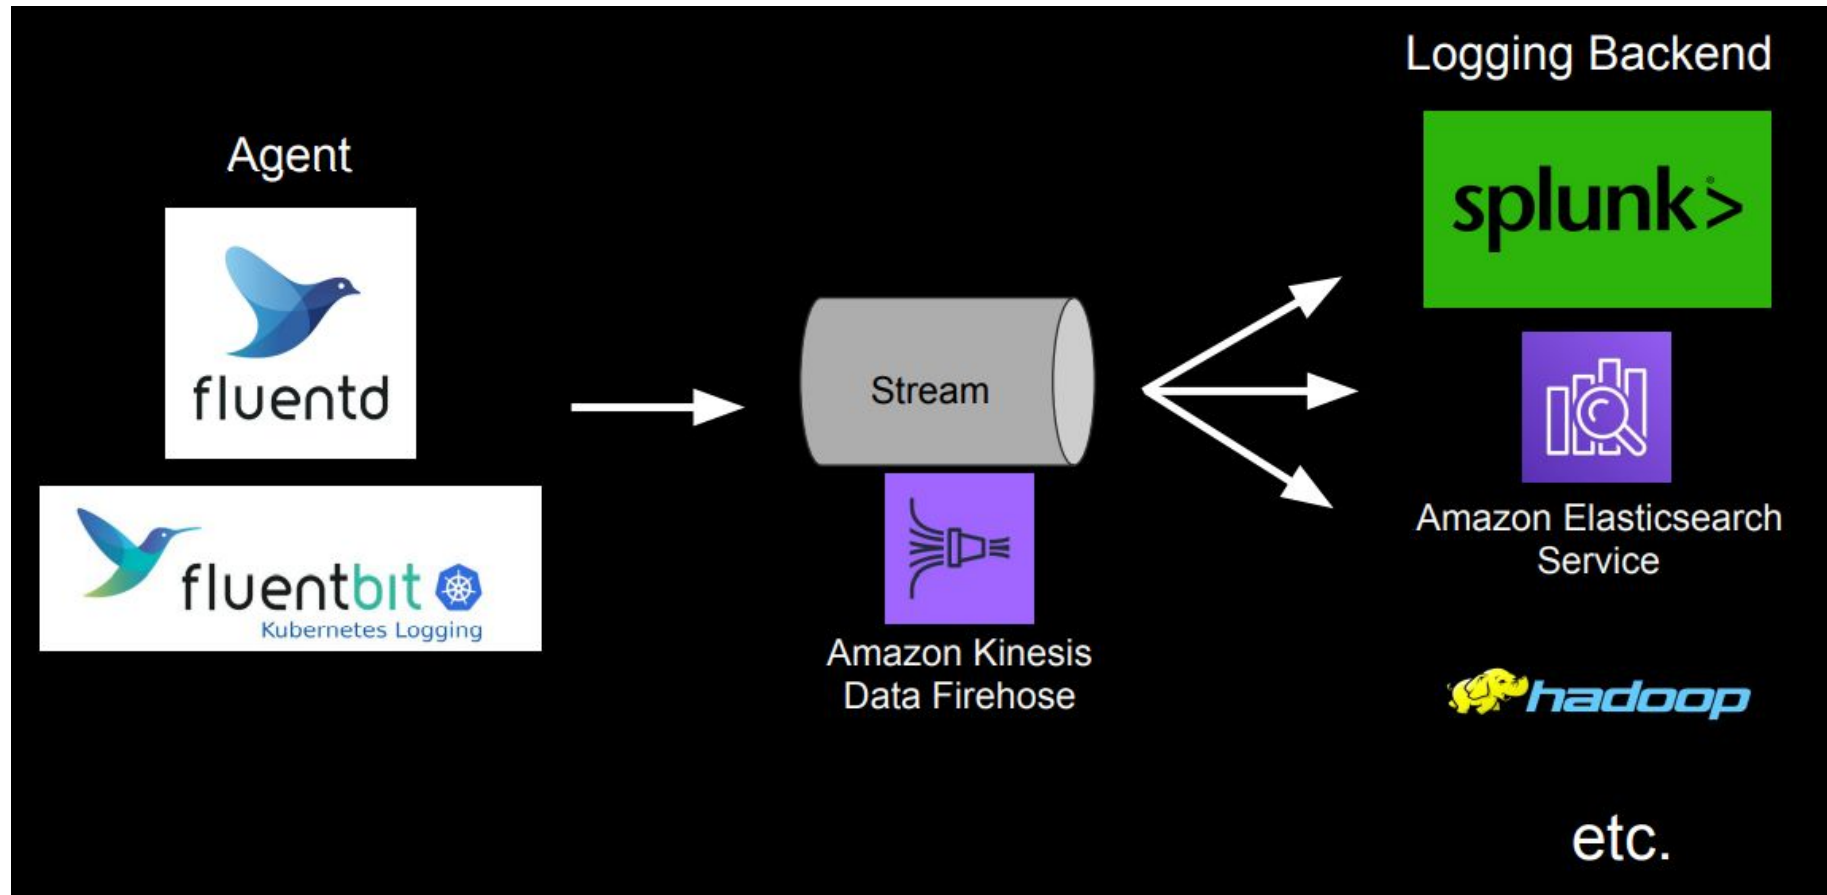


Amazon Elasticsearch
Service



etc.

For Streaming to Multiple Log Backends



EFK Stack



Amazon
Elasticsearch
Service



Under Pressure!



- fluentd has 100+ plugins, fluentbit has around 20 (April 2020)
- However as traffic goes up, fluentd can't keep up
 - fluentd based on Ruby and memory intensive
 - Slow propagation of logs
 - Loss of logs
 - fluentd buffer can be increased to solve this, but not dynamic

Contd on next slide..

Under Pressure!



- fluentbit is lightweight and keeps up with higher traffic
- Ways to solve the high traffic problem
 - fluentd to Kinesis Data Firehose to Logging Backend
 - fluentbit to Logging Backend
 - Hard to replace fluentd coz of plugin support if already existing in enterprise

Implementation

EKS Control Plane Logging


EKS Control Plane Logging

- **K8s api**
- **audit**
- **authenticator**
- **controllerManager**
- **scheduler**

EKS Control Plane Logging

- EKS Control Plane Logging

- K8s api
- audit
- authenticator
- controllerManager
- scheduler



Amazon
CloudWatch

EKS Control Plane Logging



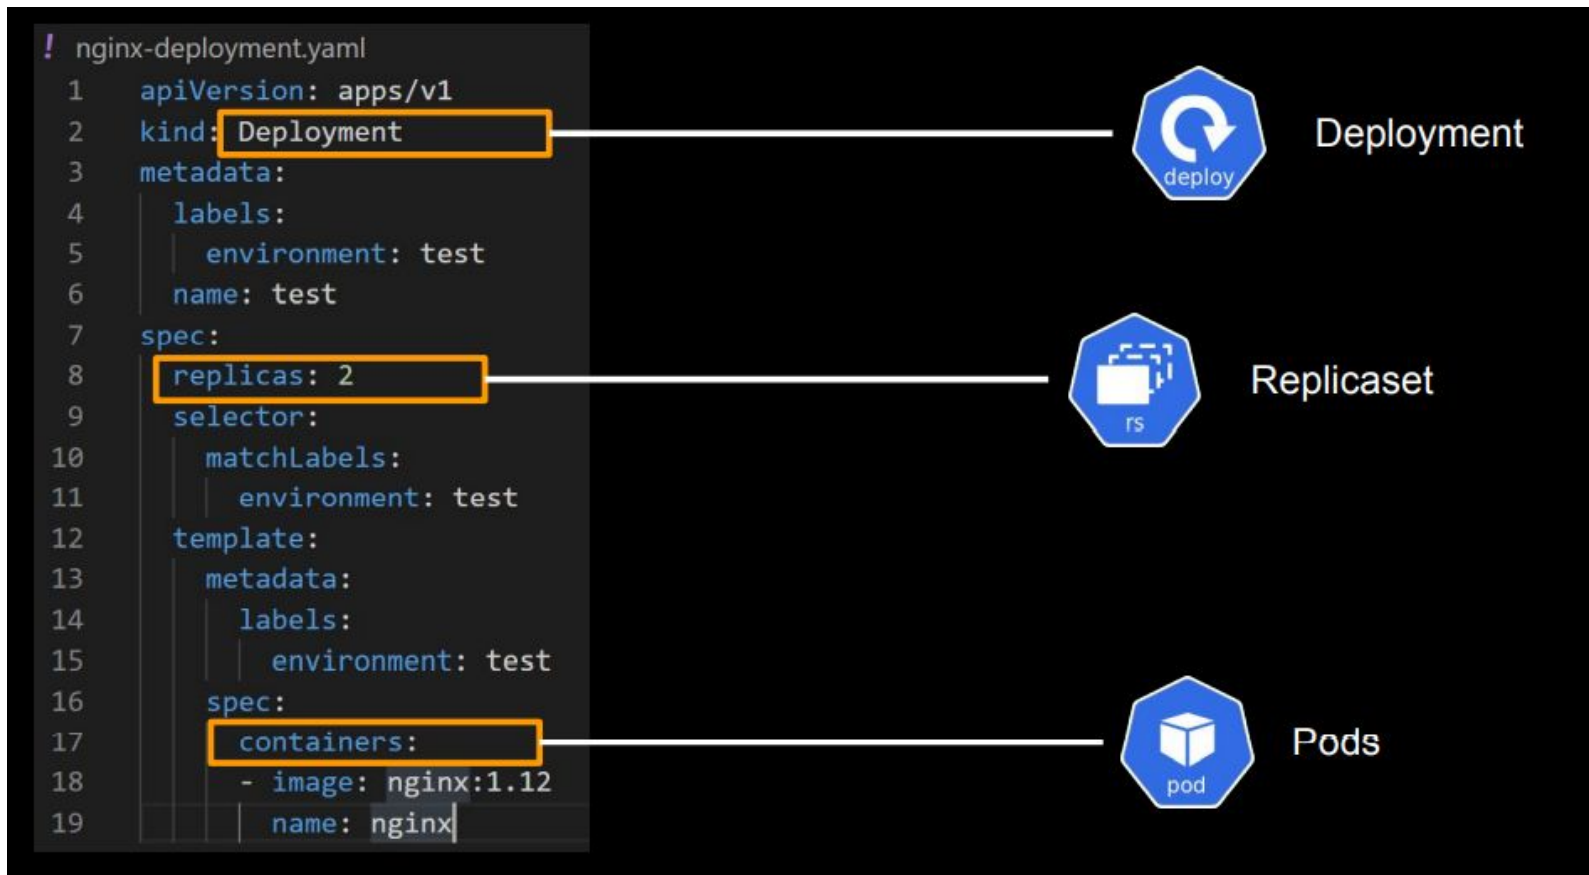
Kubernetes Dashboard

- **Web-based Kubernetes user interface**
- **Overview of applications and resources running on cluster**
- **Create and modify resources!**
 - **Pod**
 - **Deployments**
 - **Jobs**
 - **etc**

Kubernetes Dashboard

EKS ADVANCED CONCEPTS

Quiz: What Resources Created?



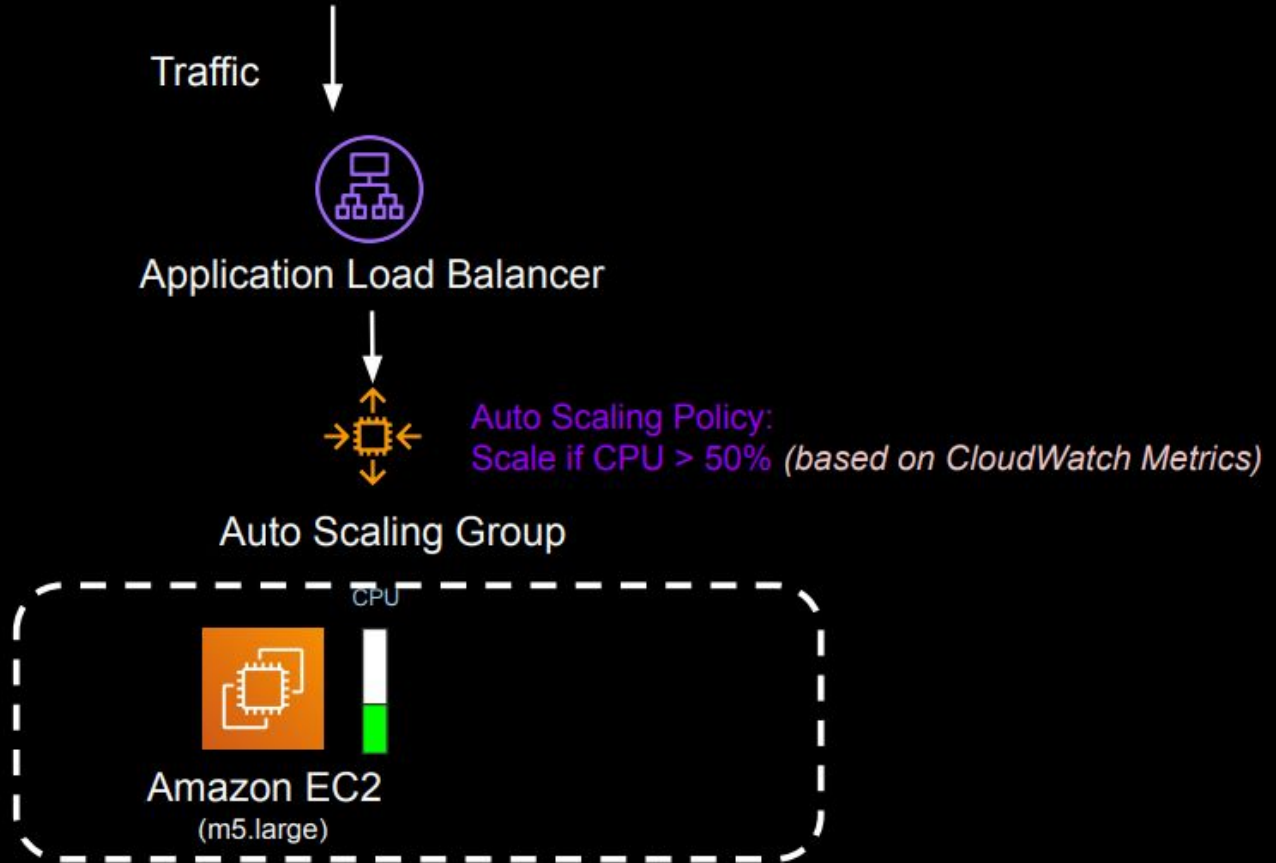
Container Scaling

- Quick Look into EC2 Scaling
- Container Scaling
- Understand Pod Limits and Requests
- Horizontal Pod Autoscaler
- Understand Manifest File
- HPA Demo

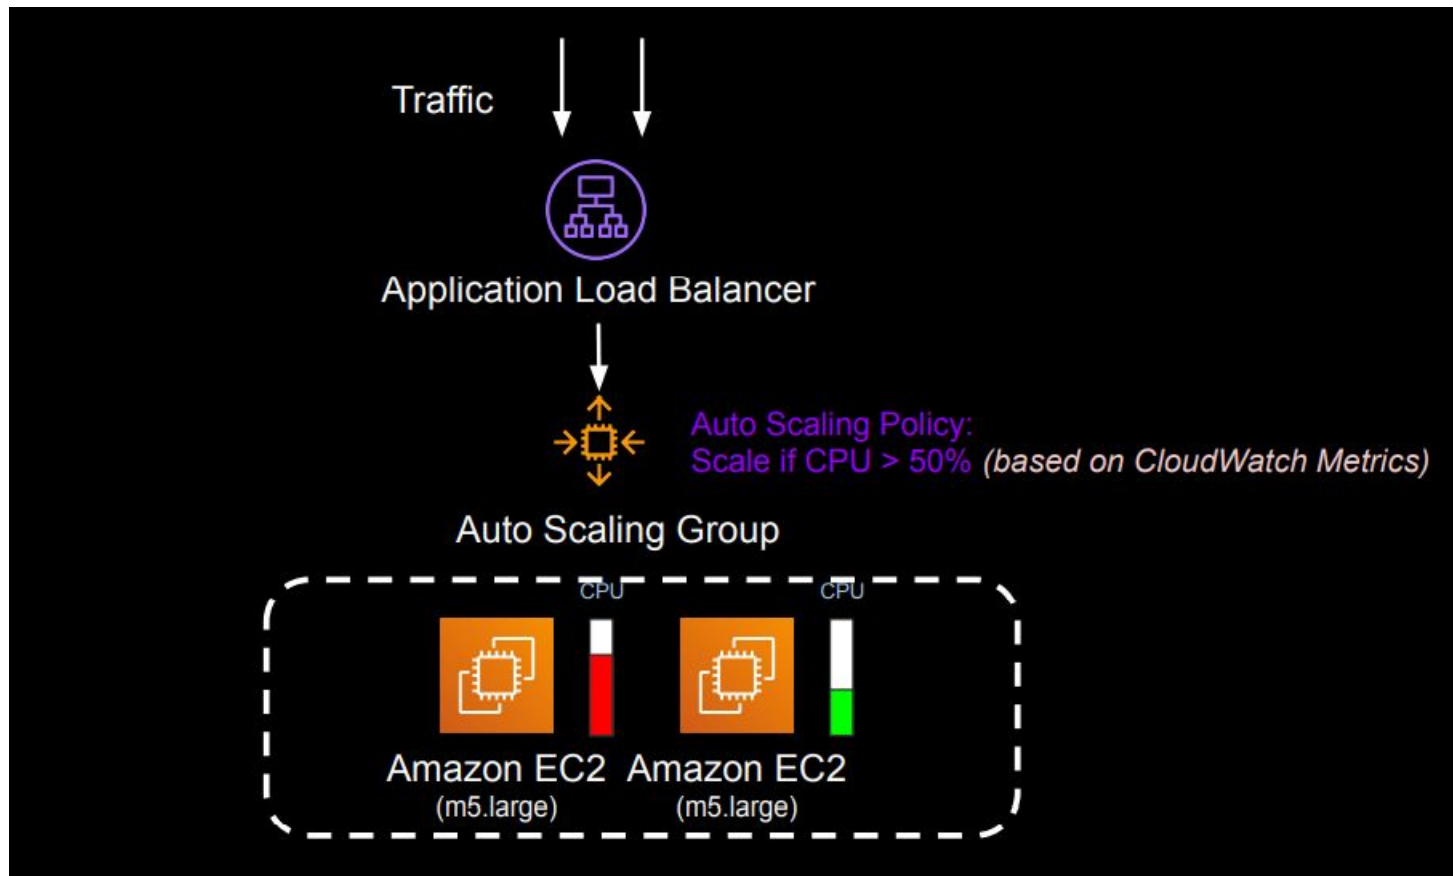
Horizontal Pod Autoscaler



Going Back to EC2 Scaling



Going Back to EC2 Scaling

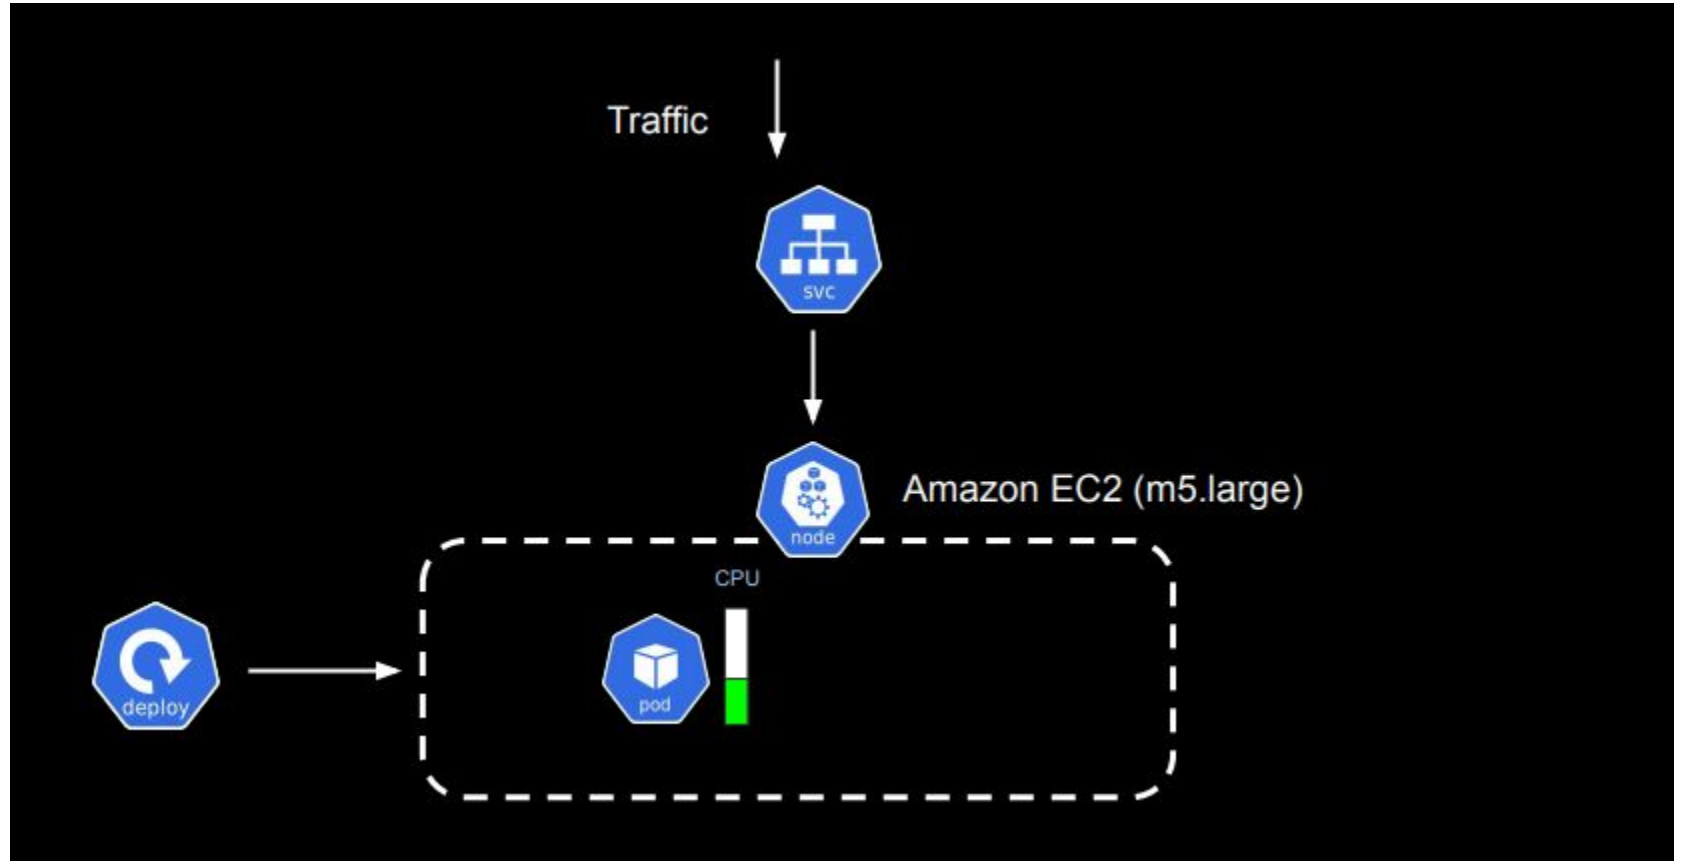


EKS Container Scaling




Amazon EC2
(m5.large)

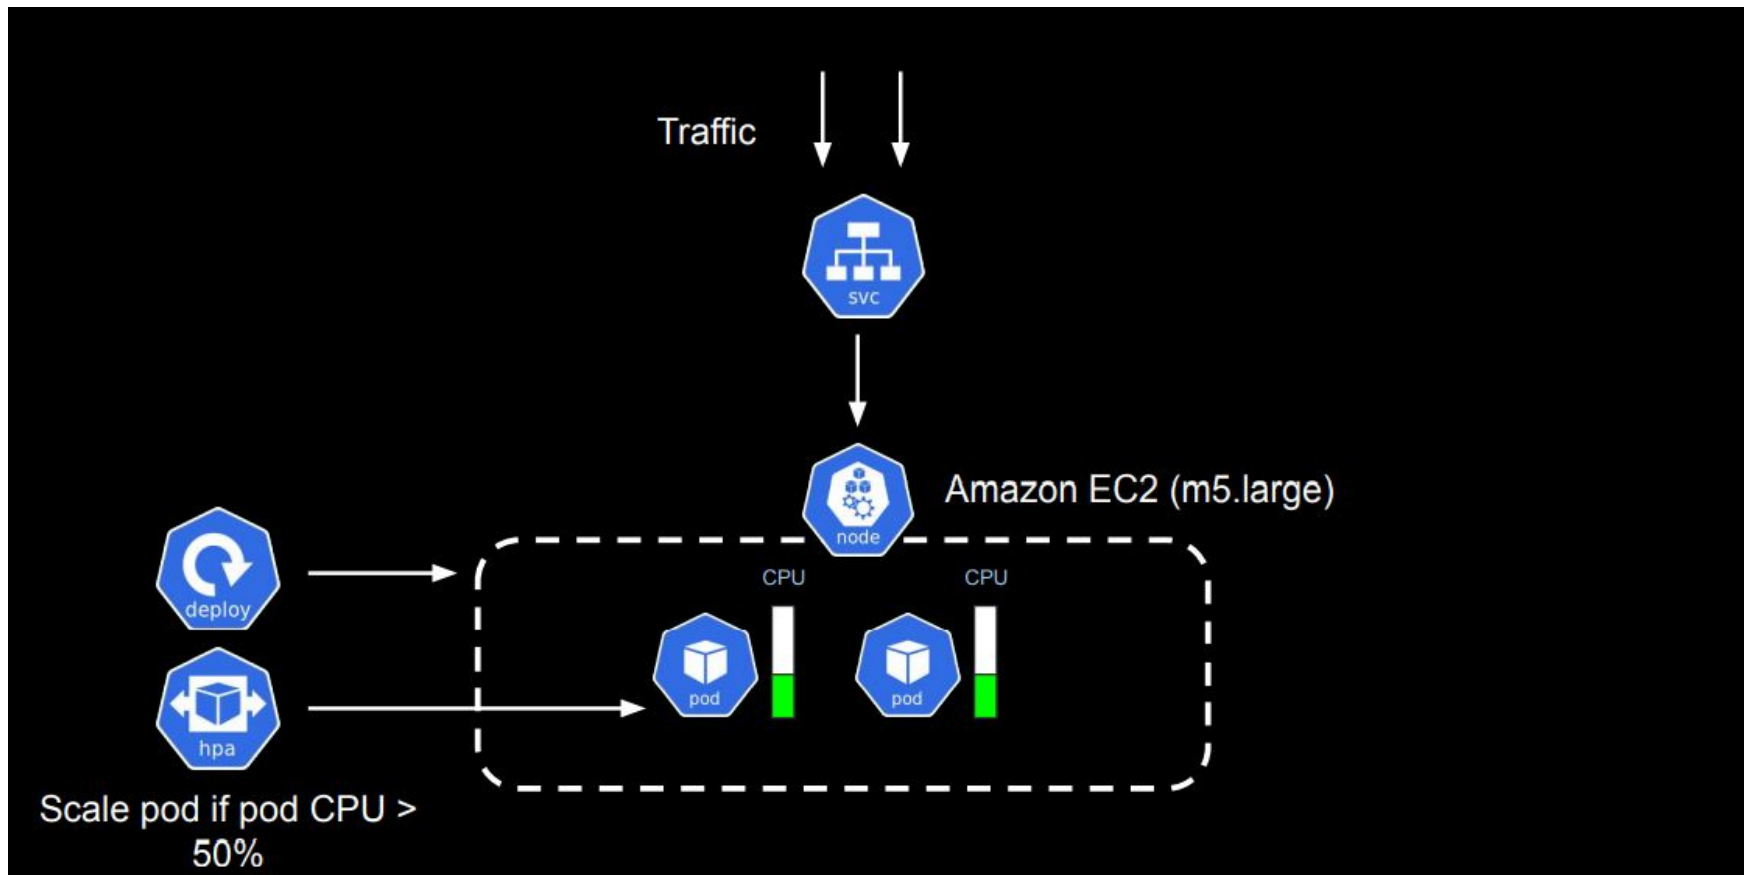
Horizontal Pod Autoscaler (HPA)



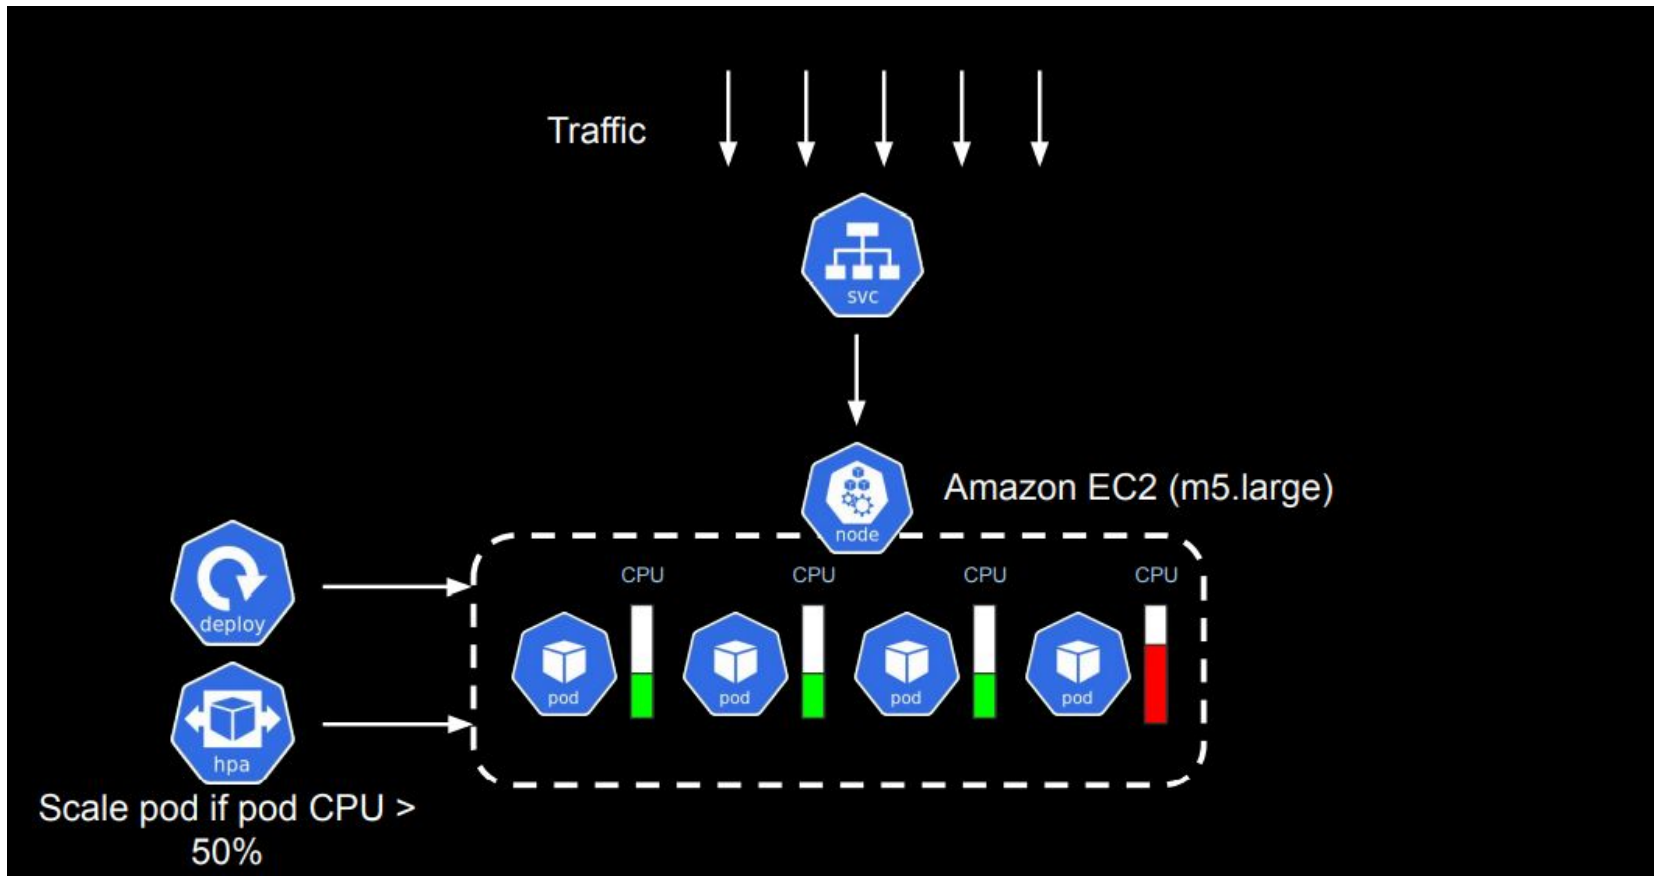
Horizontal Pod Autoscaler (HPA)



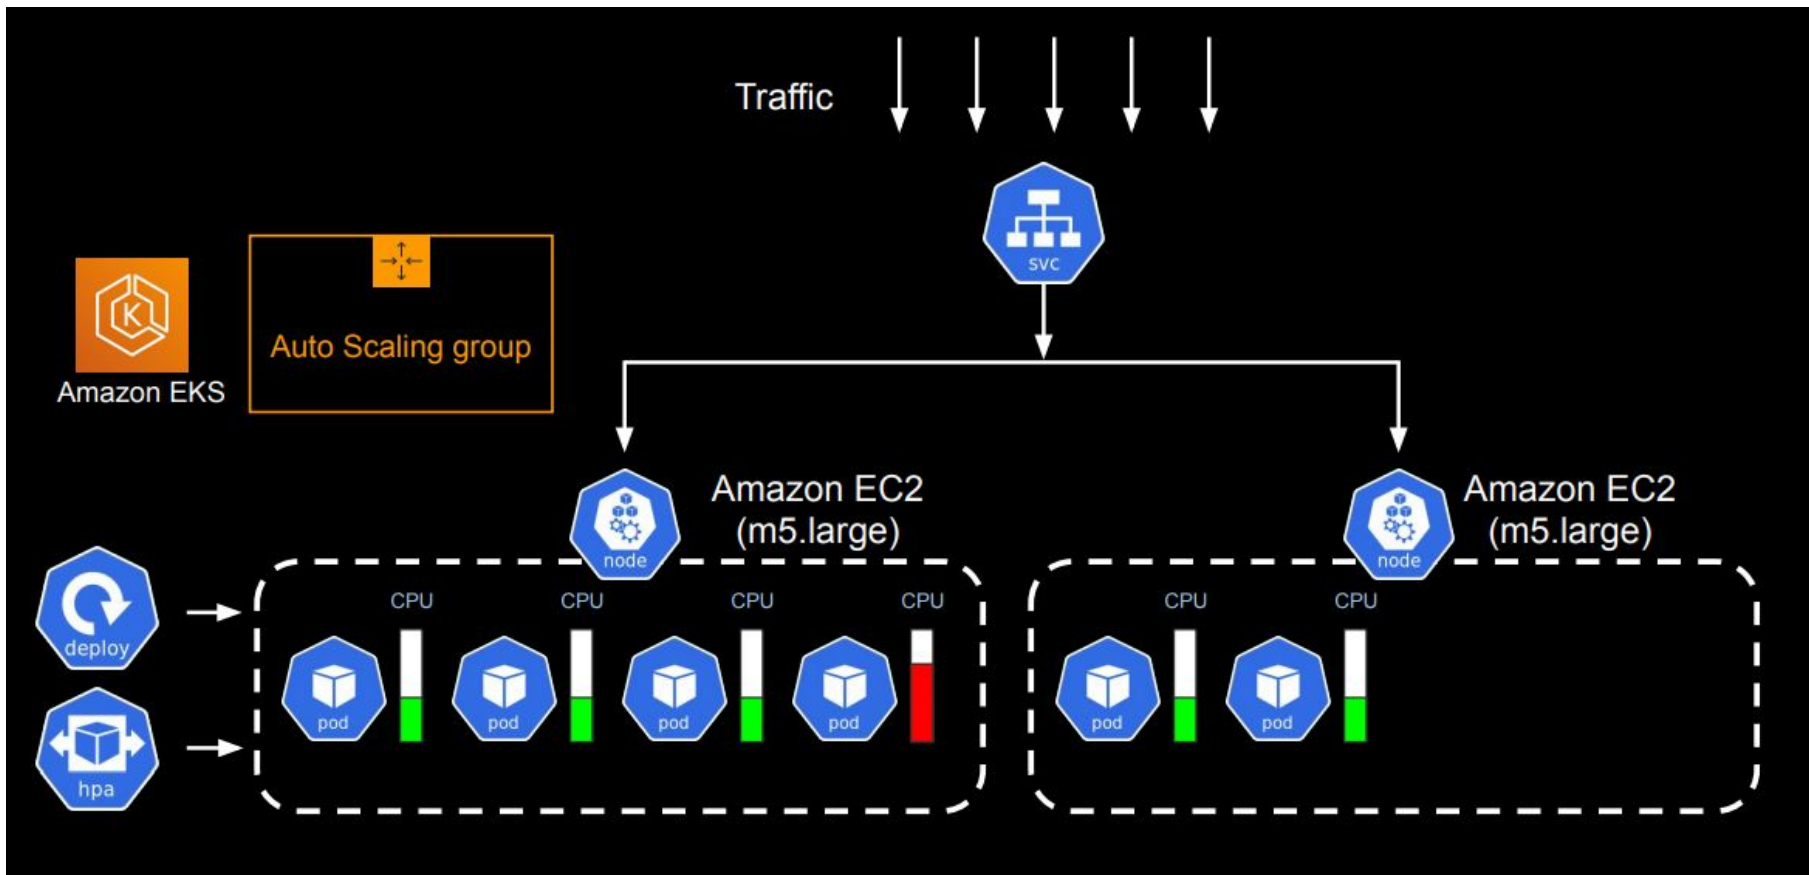
Horizontal Pod Autoscaler (HPA)




Horizontal Pod Autoscaler (HPA)



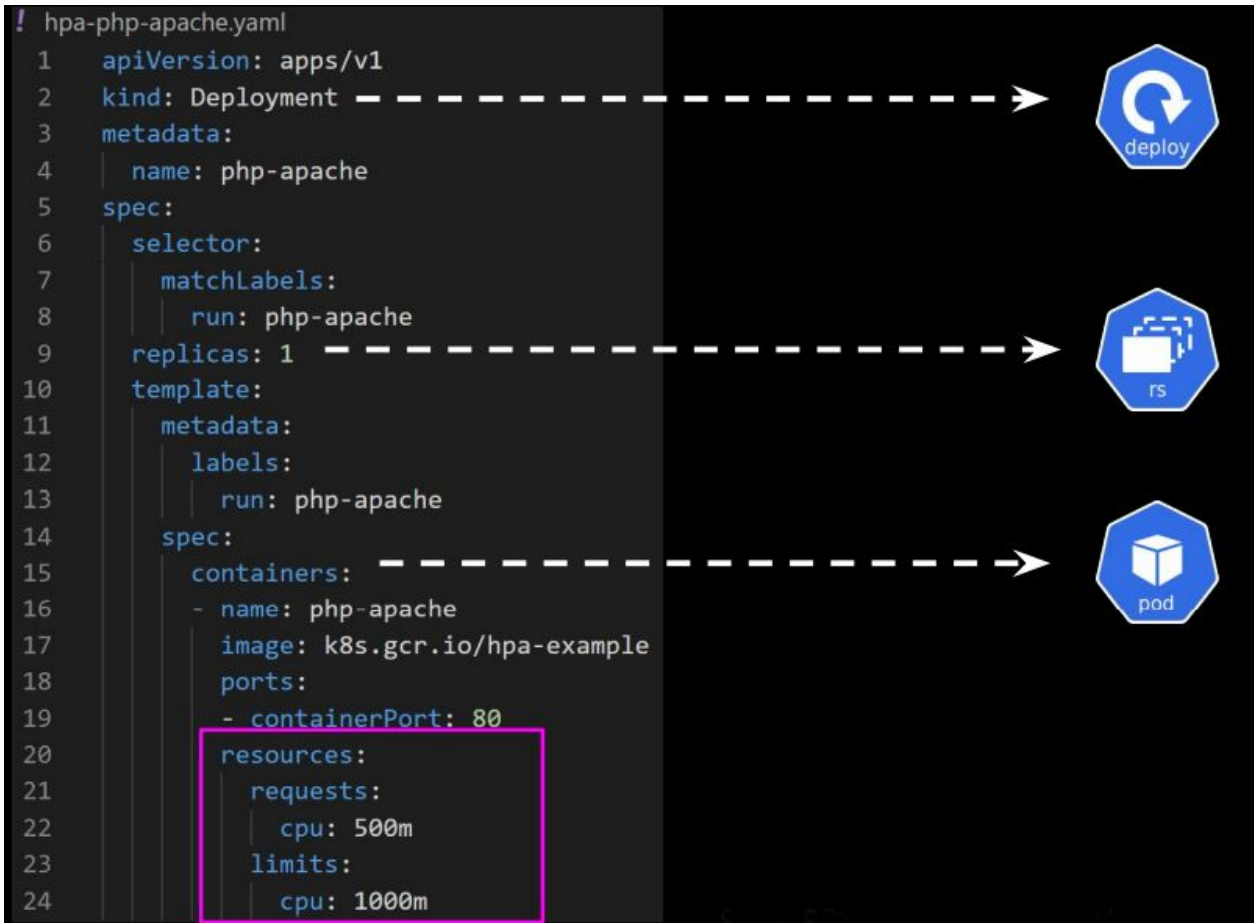
EKS Cluster Autoscaler



Horizontal Pod Autoscaler (HPA)



Pod Request And Limit



Pod Request And Limit

```
! hpa-php-apache.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: k8s.gcr.io/hpa-example
18         ports:
19         - containerPort: 80
20       resources:
21         requests:
22           cpu: 500m
23         limits:
24           cpu: 1000m
```



Amazon EC2
(m5.large)



pod

Instance Size	vCPU	Memory (GiB)
m5.large	2	8

1 vCPU = 1000m (millicore)

One pod requesting CPU of 500m (Half of 1 vCPU)

One pod CPU limit 1000m (1 vCPU)

Pod Request And Limit

```
! hpa-php-apache.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16         - name: php-apache
17           image: k8s.gcr.io/hpa-example
18           ports:
19             - containerPort: 80
20           resources:
21             requests:
22               cpu: 500m
23             limits:
24               cpu: 1000m
```

=

```
! hpa-php-apache.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16         - name: php-apache
17           image: k8s.gcr.io/hpa-example
18           ports:
19             - containerPort: 80
20           resources:
21             requests:
22               cpu: 0.5
23             limits:
24               cpu: 1
```


Pod Request And Limit

```
! hpa-cpu-and-memory-php-apache.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: k8s.gcr.io/hpa-example
18         ports:
19         - containerPort: 80
20         resources:
21           requests:
22             cpu: 500m
23             memory: 256Mi
24           limits:
25             cpu: 1000m
26             memory: 512Mi
```



Amazon EC2
(m5.large)

Instance Size	vCPU	Memory (GiB)
m5.large	2	8

1 vCPU = 1000m (millicore)

One pod requesting CPU of 500m (Half of 1 vCPU) and 256 MiB of Memory

One pod CPU limit 1000m (1 vCPU) and 512 MiB of Memory

Pod Request And Limit + HPA

! hpa-php-apache.yaml

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: k8s.gcr.io/hpa-example
18         ports:
19         - containerPort: 80
20       resources:
21         requests:
22           cpu: 500m
23         limits:
24           cpu: 1000m
```

Create Replica of Pod
If pod CPU > 50% of
request CPU
(If pod CPU exceeds
250m)



```
42  apiVersion: autoscaling/v1
43  kind: HorizontalPodAutoscaler
44  metadata:
45    name: php-apache
46    namespace: default
47  spec:
48    scaleTargetRef:
49      apiVersion: apps/v1
50      kind: Deployment
51      name: php-apache
52    minReplicas: 1
53    maxReplicas: 10
54    targetCPUUtilizationPercentage: 50
55
```

K8s HPA YAML

```
! hpa-php-apache.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: k8s.gcr.io/hpa-example
18         ports:
19         - containerPort: 80
20         resources:
21           requests:
22             cpu: 500m
23           limits:
24             cpu: 1000m
```



```
42  apiVersion: autoscaling/v1
43  kind: HorizontalPodAutoscaler
44  metadata:
45    name: php-apache
46    namespace: default
47  spec:
48    scaleTargetRef:
49      apiVersion: apps/v1
50      kind: Deployment
51      name: php-apache
52    minReplicas: 1
53    maxReplicas: 10
54    targetCPUUtilizationPercentage: 50
55
```

Pod Request And Limit

```
! hpa-php-apache.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 10
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: k8s.gcr.io/hpa-example
18         ports:
19         - containerPort: 80
20         resources:
21           requests:
22             cpu: 500m
23           limits:
24             cpu: 1000m
```

Create Replica of Pod
If pod CPU > 50% of
request CPU
(If pod CPU exceeds
250m)



```
42  apiVersion: autoscaling/v1
43  kind: HorizontalPodAutoscaler
44  metadata:
45    name: php-apache
46    namespace: default
47  spec:
48    scaleTargetRef:
49      apiVersion: apps/v1
50      kind: Deployment
51      name: php-apache
52    minReplicas: 1
53    maxReplicas: 10
54    targetCPUUtilizationPercentage: 50
55
```


HPA Demo on EKS

- Install Metrics Server
- Deploy Deployment, Service, HPA
- Increase Load
- Pod scale!

Note - Need node larger than t3.micro

<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough>

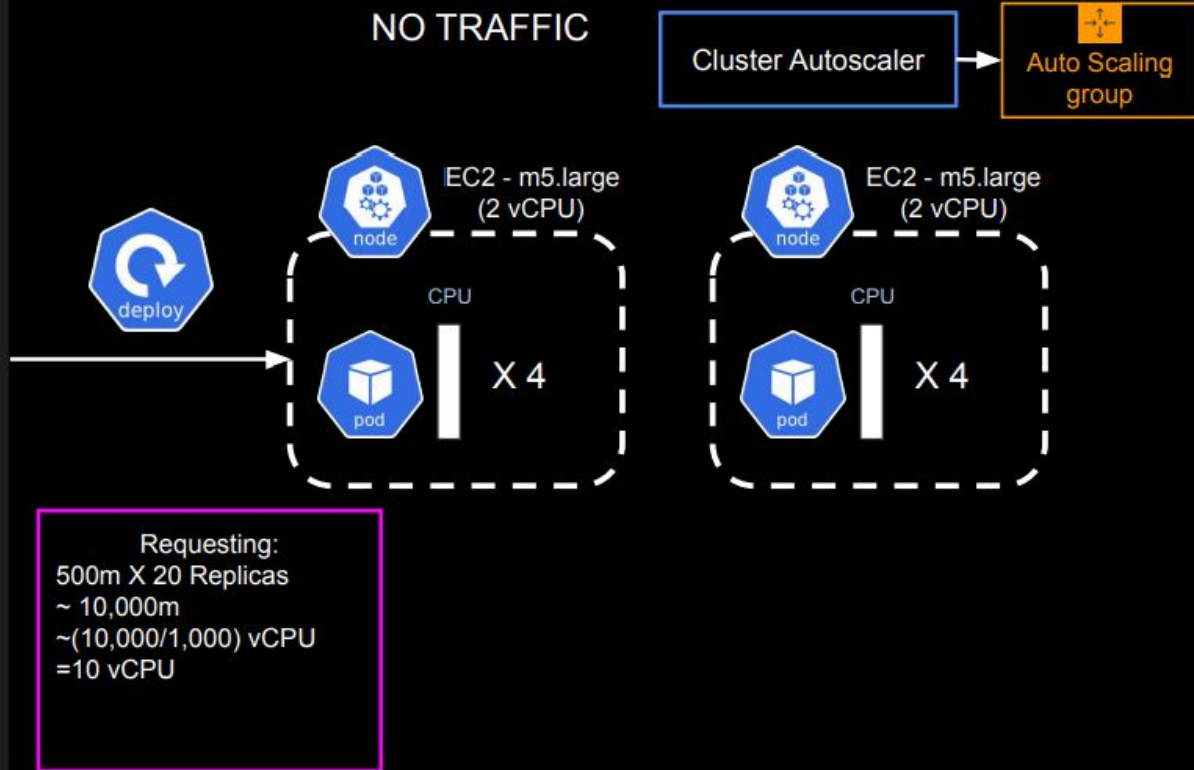
Horizontal Pod Autoscaler (HPA)



Cluster Autoscaler

```
! cluster-autoscaler-deployment-1.yaml
```

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 20
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: k8s.gcr.io/hpa-example
18         ports:
19         - containerPort: 80
20       resources:
21         requests:
22           cpu: 500m
23         limits:
24           cpu: 1000m
```



CLUSTER AUTOSCALER

! cluster-autoscaler-deployment-1.yaml

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 20
10  template:
11    metadata:
12      labels:
13        run: php-apache
14    spec:
15      containers:
16      - name: php-apache
17        image: k8s.gcr.io/hpa-example
18        ports:
19        - containerPort: 80
20      resources:
21        requests:
22          cpu: 500m
23        limits:
24          cpu: 1000m
```



CPU



Vertical Pod Autoscaler



Do NOT use this in Production!


- Restarts PODs

Vertical Pod Autoscaler

- Vertical Vs Horizontal Scaling
 - Vertical - Going up in Size
 - Horizontal - Create more of same size
- Pods will go up or down in size (after restart!)
- Used in dev to determine optimal CPU and memory for the Pod
 - VPA recommends pod request, limit
 - Use the numbers for request, limits, HPA
 - VPA should not be used with HPA
- Accept VPA recommendation with grain of salt!

EKS Ingress

Ingress - What and Why?




```
! loadbalancer-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: lb-service
5    labels:
6      app: lb-service
7  spec:
8    type: LoadBalancer
9    ports:
10     - port: 80
11    selector:
12      app: frontend
```

Ingress - What and Why?


I want to run Bat Cave Systems using Kubernetes




http://track-joker.com



http://monitor-batcave.com



http://order-new-batsuit.com




Ingress - What and Why?

- Costly to maintain one Load Balancer for each service
 - No URL based routing
 - Maintenance overhead of managing all separate services
- "Alfred HELP ME"




<https://batcave.com/track-joker>
<https://batcave.com/monitor-batcave>
<https://batcave.com/order-new-batsuit>


<http://track-joker.com>



<http://monitor-batcave.com>



<http://order-new-batsuit.com>



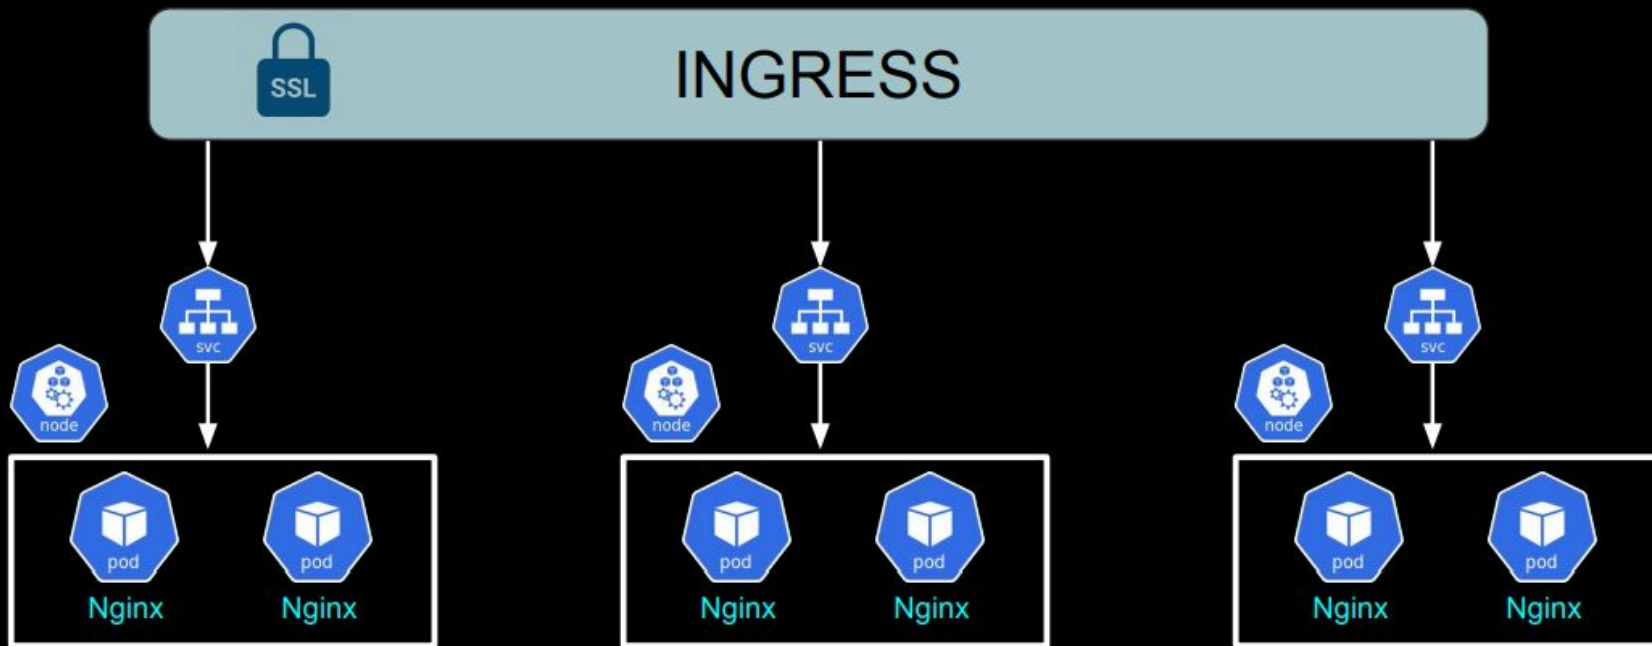
Ingress - What and Why?

Master Bruce, you need Ingress!



Ingress - What and Why?

<https://batcave.com/track-joker>
<https://batcave.com/monitor-batcave>
<https://batcave.com/order-new-batsuit>



Ingress Controllers

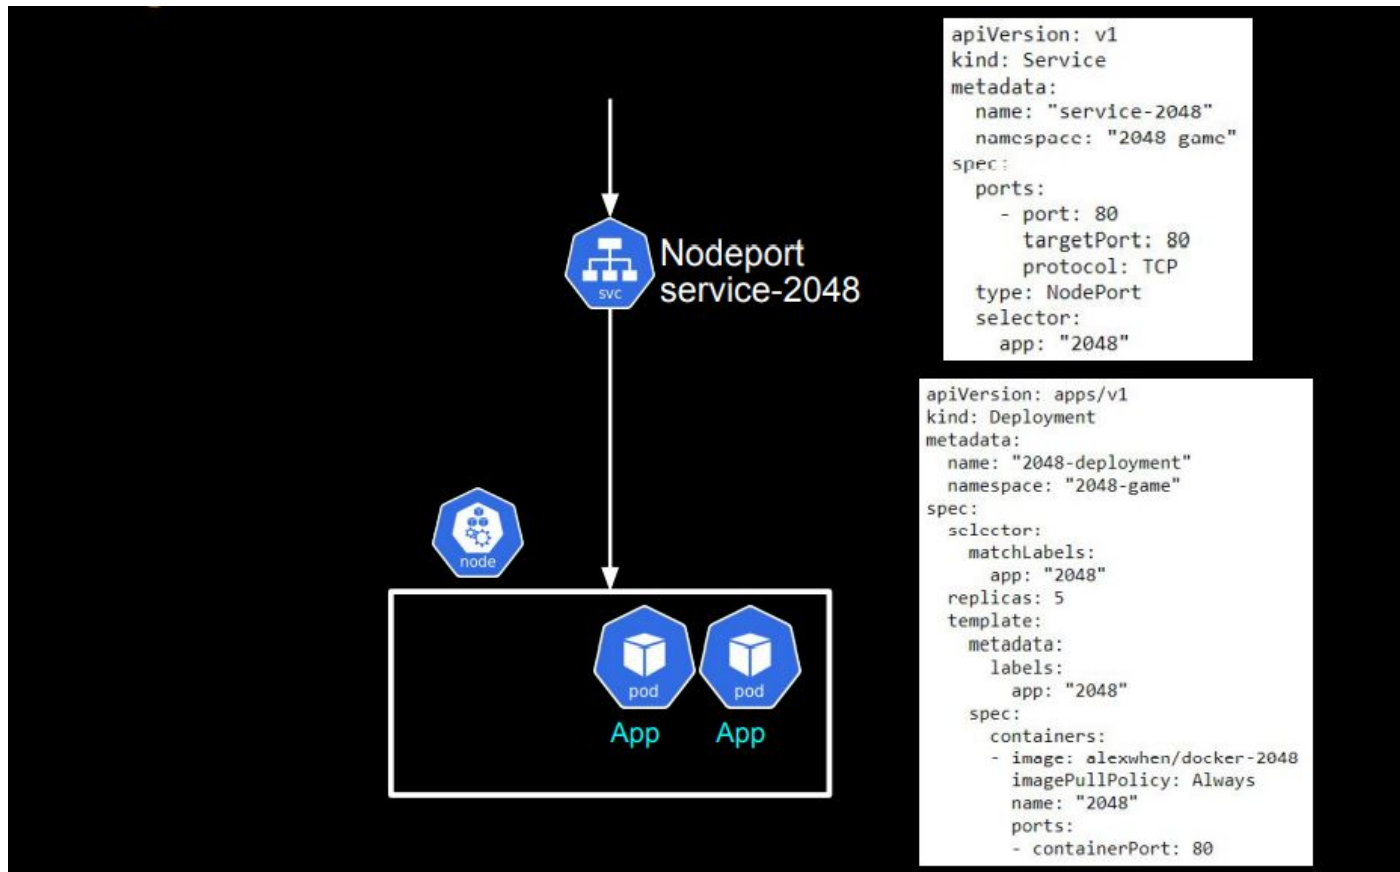


HAPROXY

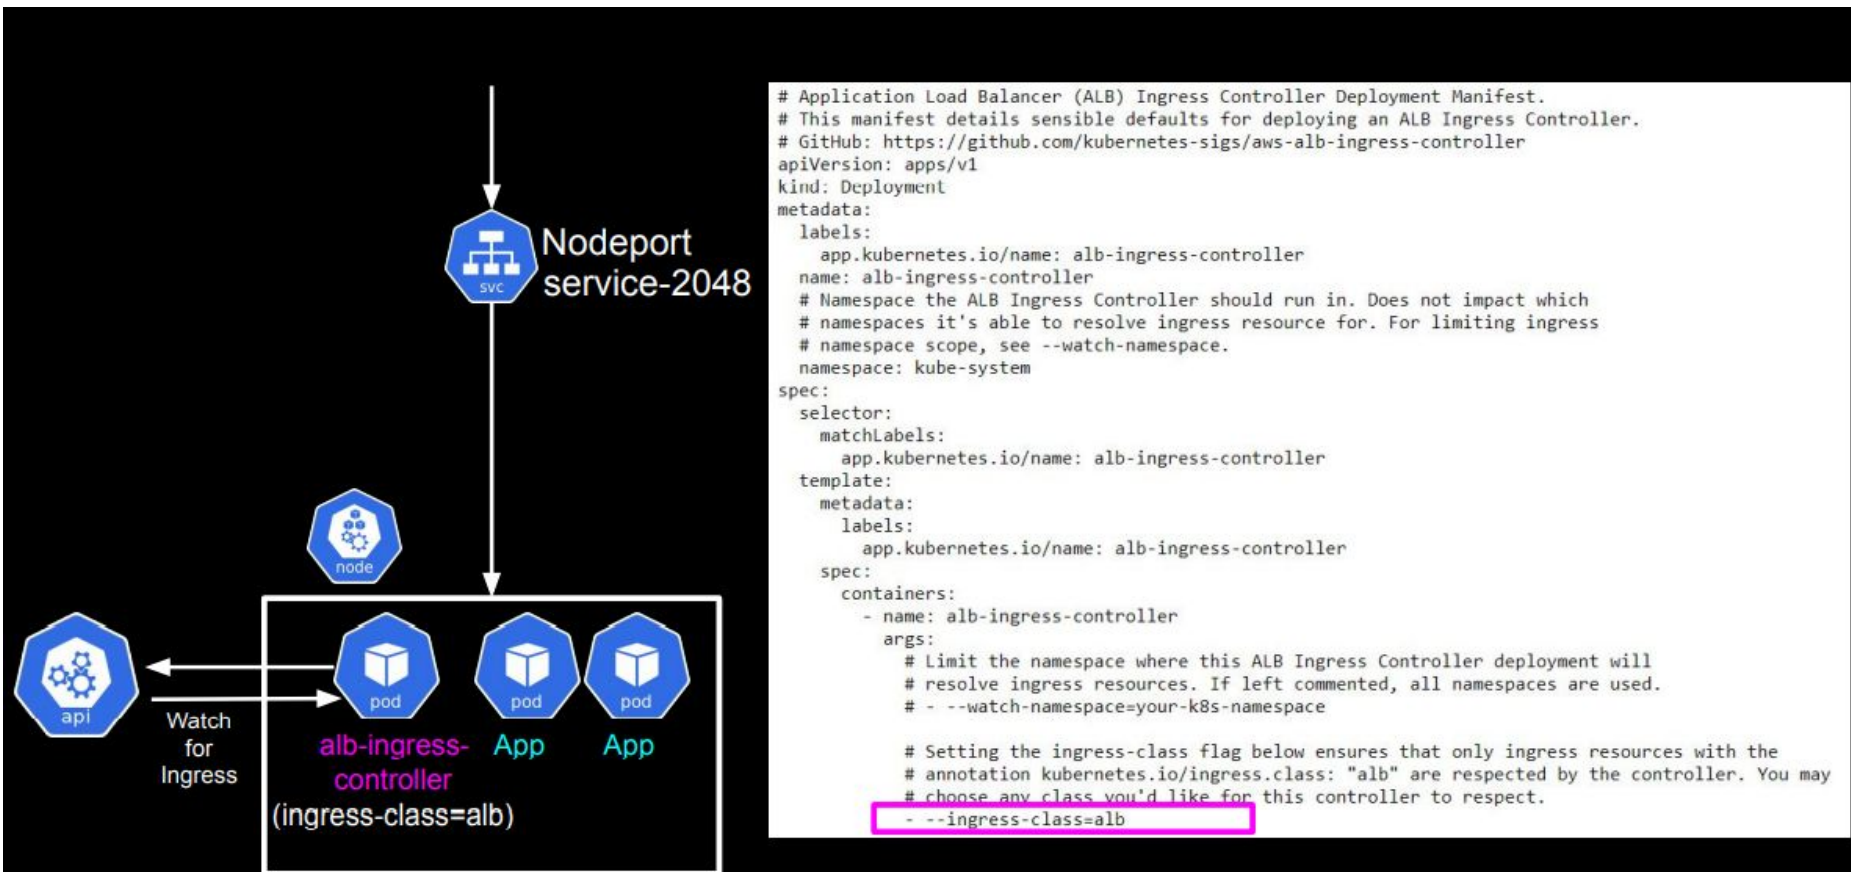


alb-ingress-controller
+
AWS ALB

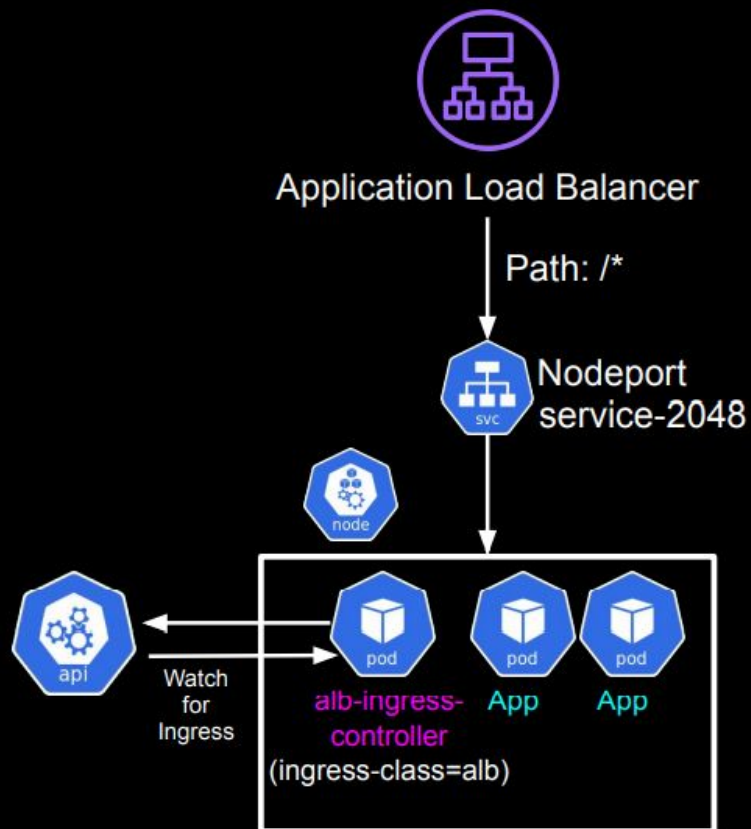
ALB Ingress



ALB Ingress Controller

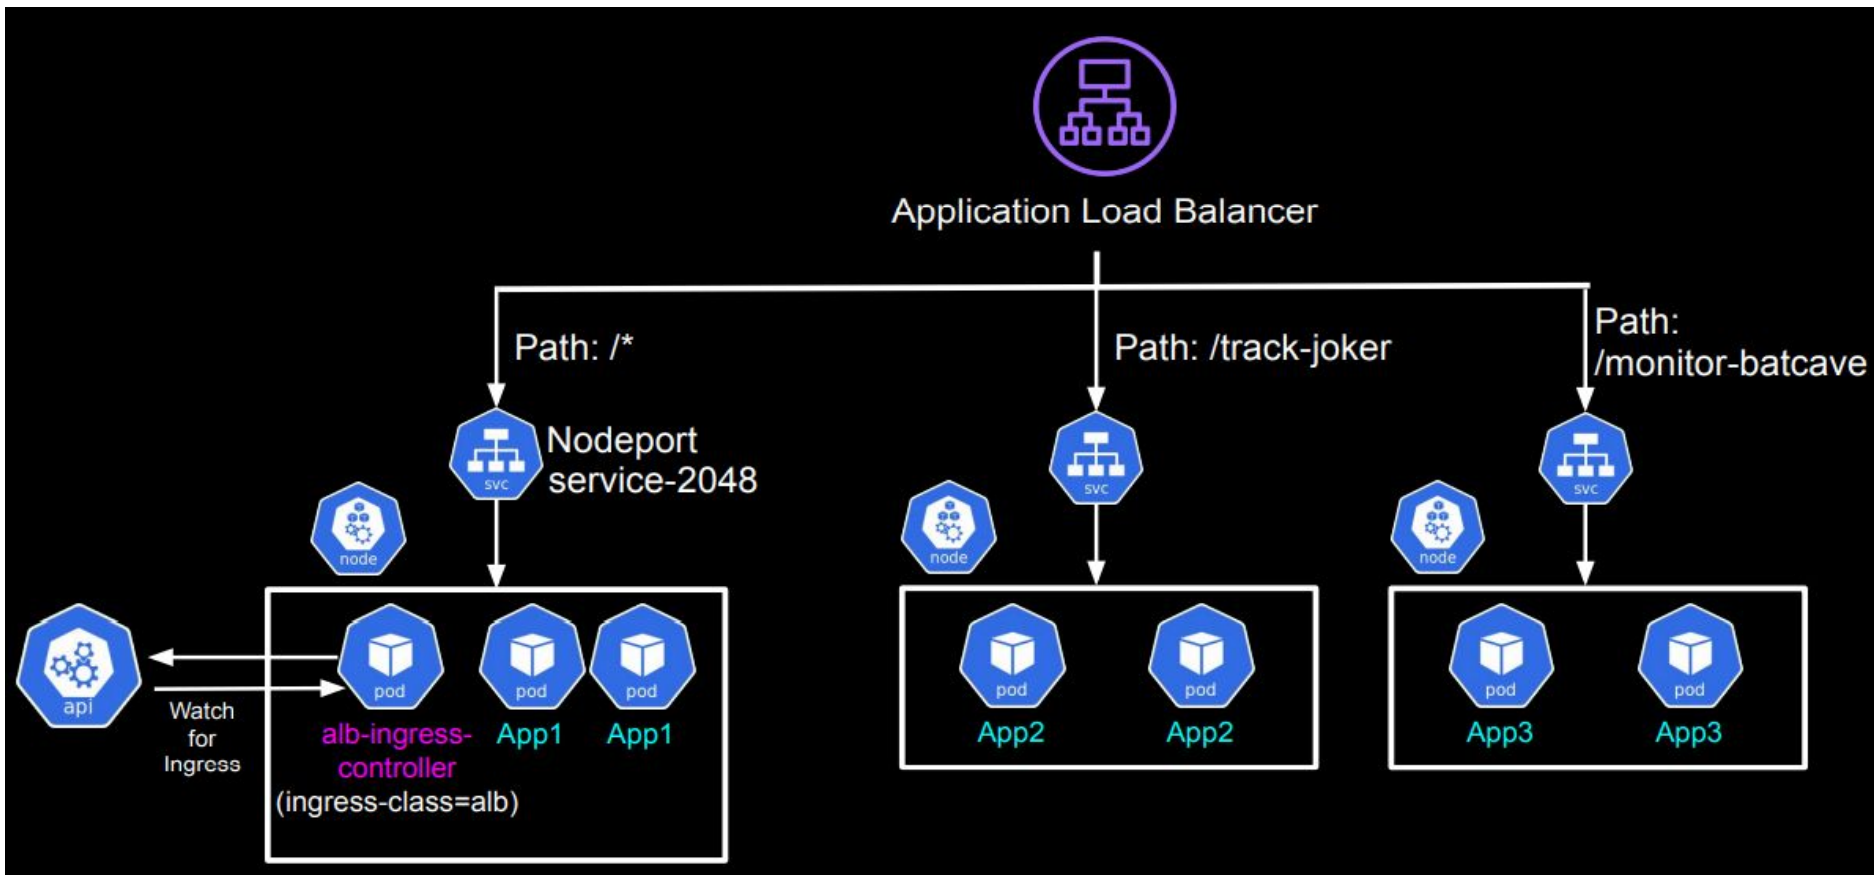


Ingress resource



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: "2048-ingress"
  namespace: "2048-game"
  annotations:
    kubernetes.io/ingress.class: alb
    alb.ingress.kubernetes.io/scheme: internet-facing
  labels:
    app: 2048-ingress
spec:
  rules:
    - http:
        paths:
          - path: /*
            backend:
              serviceName: "service-2048"
              servicePort: 80
```

ALB Ingress



Ingress - What and Why?

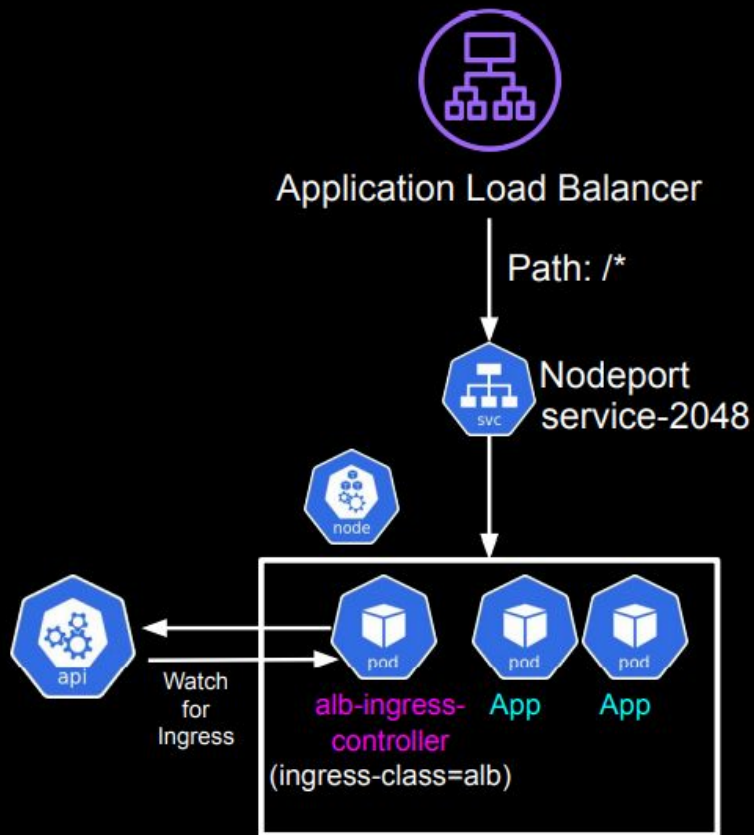
INGRESS CONTROLLER

- Monitors Ingress resources
- Creates necessary AWS resources for Ingress
 - Such as ALB for ALB Ingress Controller
- One Cluster can have more than one Ingress Controller!
 - Ingress Resource defines which Ingress Controller to use

INGRESS RESOURCE

- Selects which Ingress Controller to use
- Defines the URL Path and corresponding backend Service

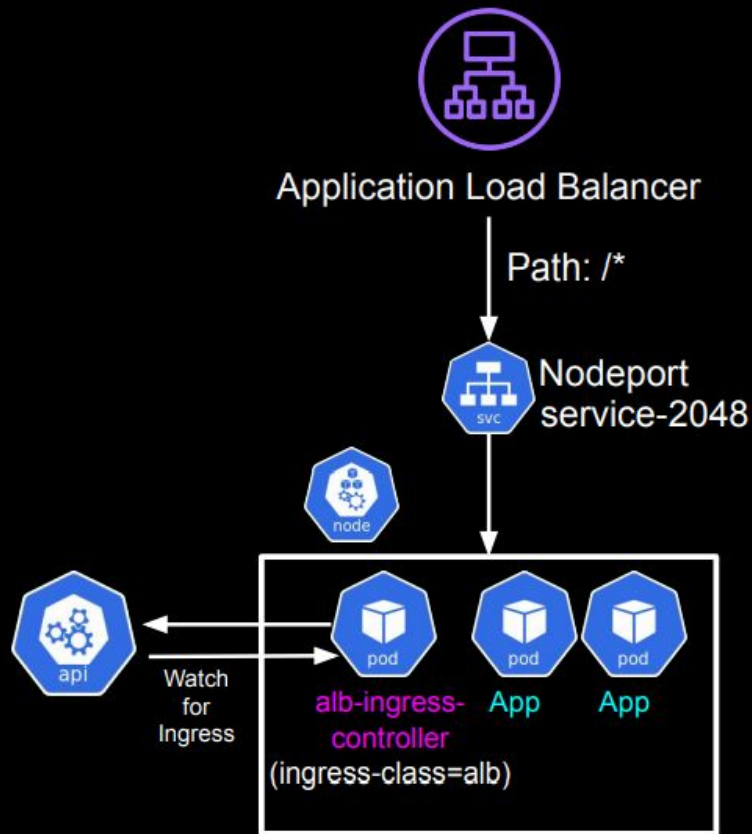
ALB Ingress



INGRESS CONTROLLER

- Creates ALB and AWS Resources
 - Directed by Ingress Resource
- Require proper IAM Policy
- Require Service Account and IAM Role for the alb-ingress-controller pod

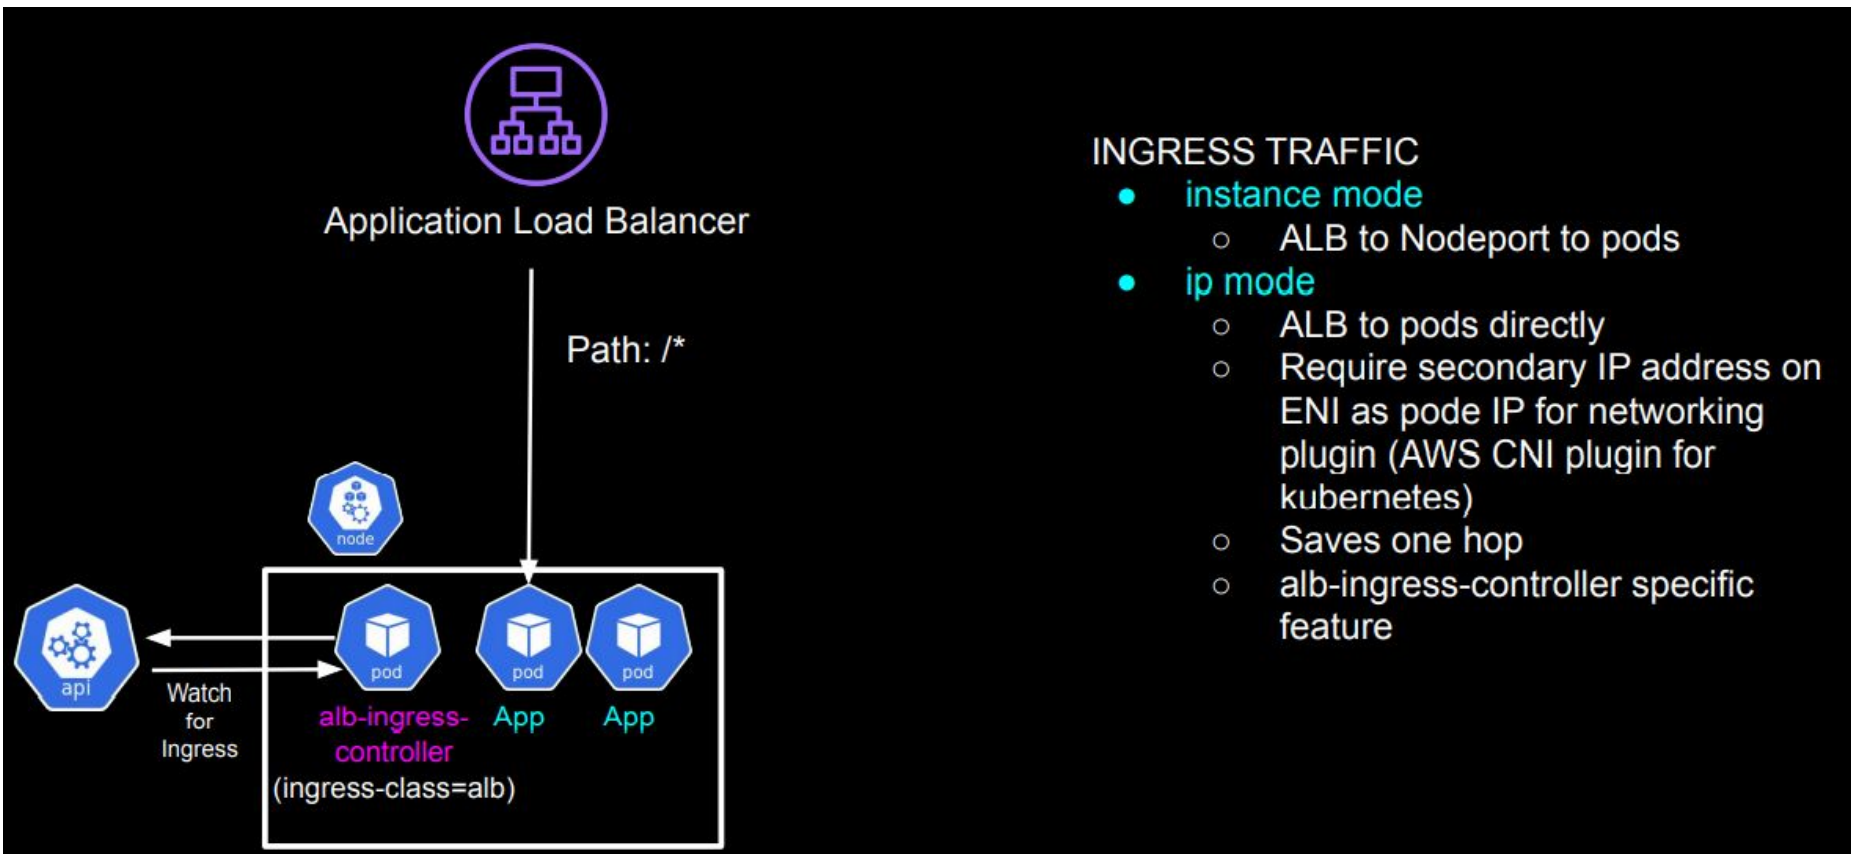
ALB Ingress



INGRESS TRAFFIC

- **instance mode**
 - ALB to Nodeport to pods
- **ip mode**
 - ALB to pods directly
 - Require secondary IP address on ENI as pod IP for networking plugin (AWS CNI plugin for kubernetes)
 - Saves one hop
 - alb-ingress-controller specific feature

ALB Ingress



Thank You